

Solving Partial Differential Equations (PDEs) with Quantum Computers

Zhixin (Jack) Song, Spencer H. Bryngelson

Georgia Institute of Technology

09/22 IEEE Quantum Week 2023, Bellevue, WA

Speakers




Zhixin (Jack) Song
3rd-year Ph.D. Student
zsong300@gatech.edu



Prof. Spencer Bryngelson
PI, shb@gatech.edu

What we do



Computational Fluid Dynamics (CFD)



Quantum Computing (QC)

Timeline


- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)

Timeline

- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)


What is a PDE?

An equation to solve with partial derivatives.




Wave equation

$$\frac{\partial^2 U}{\partial t^2} = c^2 \nabla^2 U$$



Schrödinger equation

$$i\hbar \frac{\partial}{\partial t} \psi = -\frac{\hbar^2}{2m} \nabla^2 \psi + mV\psi$$




Navier-Stokes equations

$$\frac{\partial \mathbf{v}}{\partial t} + (\mathbf{v} \cdot \nabla) \mathbf{v} = \nu \nabla^2 \mathbf{v} - \nabla P$$


Why quantum computing?

Solving such problems in large scale is expensive! Each high resolution CFD simulation could cost millions of dollars and produce ~tons of CO₂.



Example CFD simulation using
5.3 billion nodes (~32 qubits)

Why quantum computing?



Apple's 5nm process M1 Ultra chip contains 114 billion transistors. As the size of transistors gradually shrinks to a few atoms, it will no longer be a reliable bit due to quantum tunneling.

Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World
in Data

Transistor count

50,000,000,000

10,000,000,000

5,000,000,000

1,000,000,000

500,000,000

100,000,000

50,000,000

10,000,000

5,000,000

1,000,000

500,000

100,000

50,000

10,000

5,000

1,000

500

100

50

10

5

1

Data source: Wikipedia (wikipedia.org/wiki/Transistor_count)

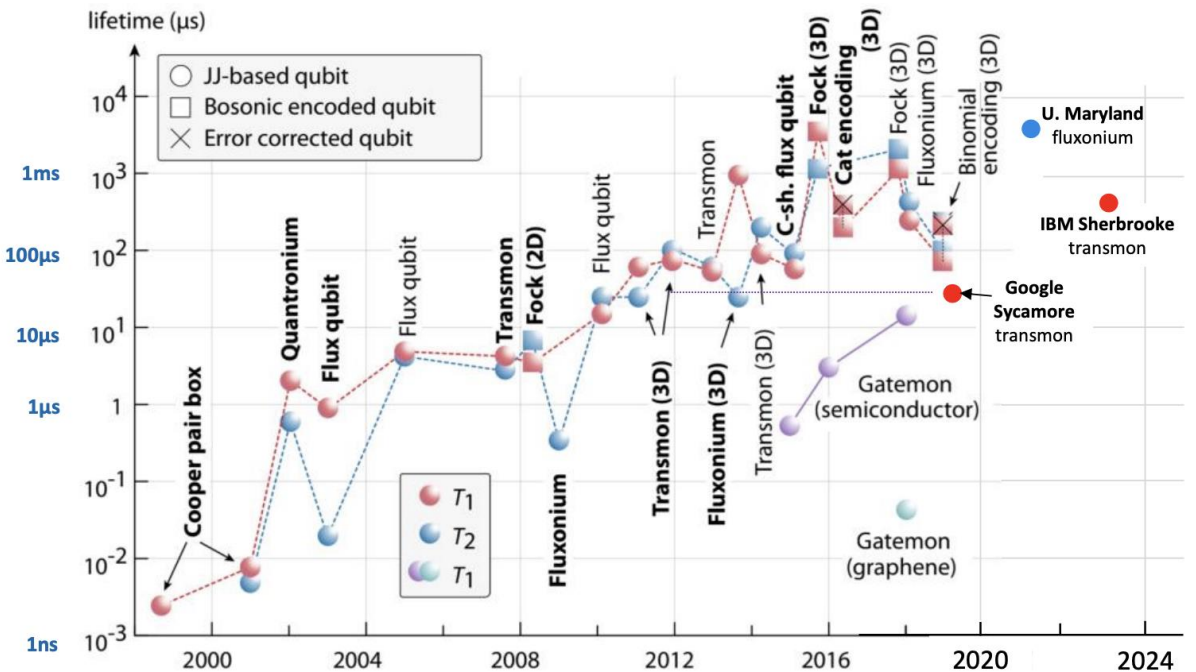

OurWorldinData.org – Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.

Year in which the microchip was first introduced

Why quantum computing?

- Computational space doubles every time you add a qubit



Why quantum computing?

- No heat created (in principle) according to Landauer's principle
- Energy consumption is several orders of magnitude lower

We estimate the total average power consumption of our apparatus under worst-case conditions for chilled water production to be **26 kW**. This power does not change appreciably between idle and running states of the quantum processor, and it is also independent of the circuit depth. This means that the energy consumed during the 200 s required to acquire 1M samples in our experiment is $\sim 5 \times 10^6$ J (~ 1 kWh). As compared to the qFlex classical simulation on Summit, we require roughly 7 orders of magnitude less energy to perform the same computation (see Table VII). Furthermore, the data acquisition time is currently dominated by control hardware communications, leading to a quantum processor duty cycle as low as 2%. This means there is significant potential to increase our energy efficiency further.

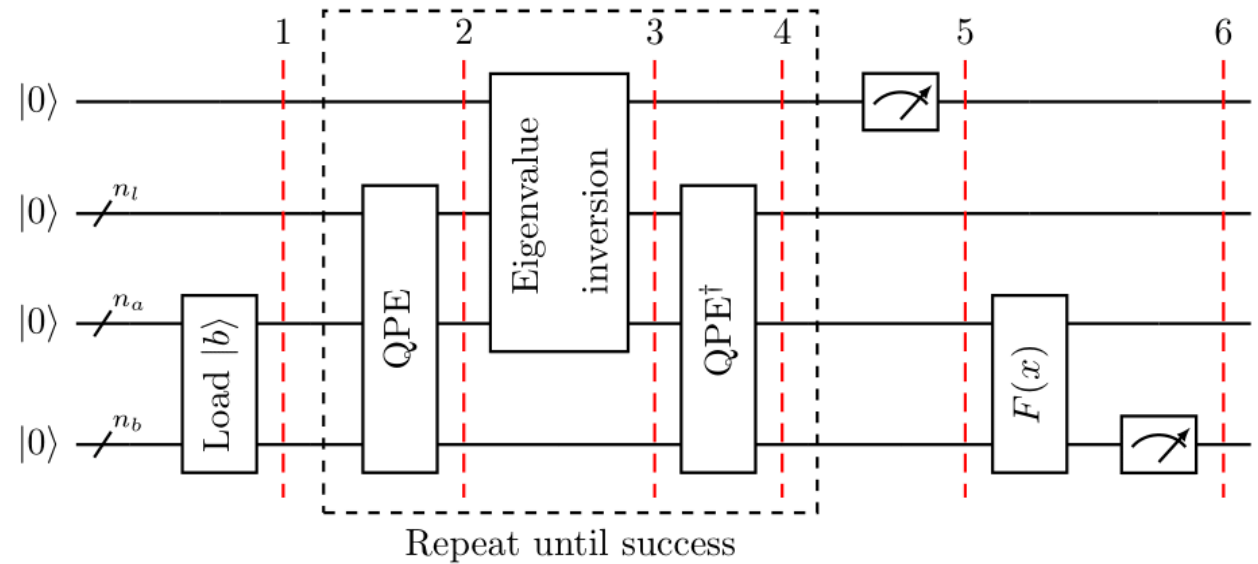

Year	Computer	Power consumption (KW)
2022	Frontier – Oak Ridge National Lab, USA	21,100
2020-2021	Fugaku - RIKEN Center for Computational Science, Japan	29,889
2018-2019	Summit – Oak Ridge National Lab, USA	10,096
2016-2017	Sunway TaihuLight – National Supercomputing Center in Wuxi, China	15,371
2013-2015	Tianhe-2A - National Supercomputer Center in Guangzhou, China	17,808
2012	Titan - Oak Ridge National Lab, USA	8,209
2011	K computer, RIKEN Advanced Institute for Computational Science, Japan	12,660
2010	Tianhe-1A - National Supercomputing Center in Tianjin, China	4,040
2009	Jaguar - Oak Ridge National Lab, USA	6,950
2008	Road Runner – Los Alamos National Lab, USA	2,483
2004-2007	<u>BlueGene/L</u> - DOE/NNNSA/LLNL, USA	2,329
2002-2003	Earth-Simulator, Japan Agency for Marine-Earth Science and Technology, Japan	3,200

<https://insidehpc.com/2023/06/the-energy-advantage-of-quantum-computers/>

Liu, J., Jiang, H., & Shen, Z. J. M. (2023). Potential Energy Advantage of Quantum Economy. *arXiv preprint arXiv:2308.08025*.

Why quantum computing?

- The core of fast CFD simulation is a large sparse linear solver
- Efficient quantum algorithms exist for solving linear systems

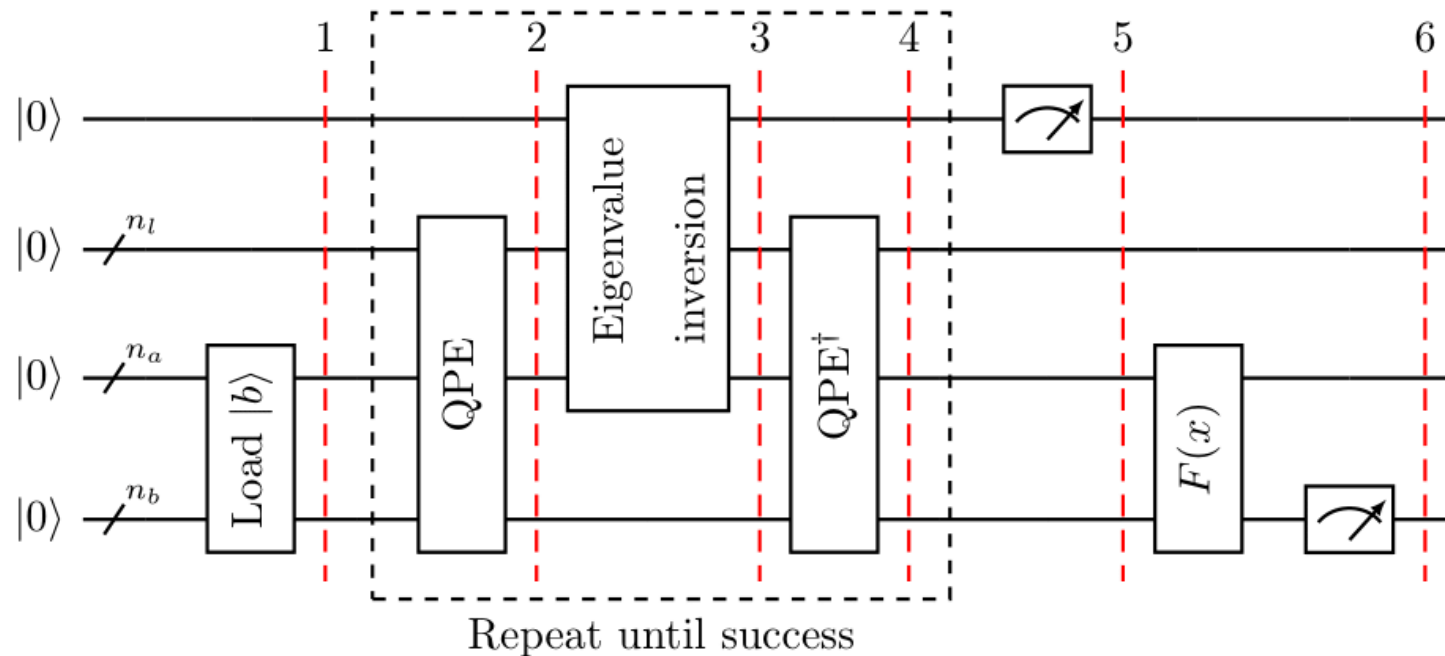


https://developer.apple.com/documentation/accelerate/sparse_solvers/

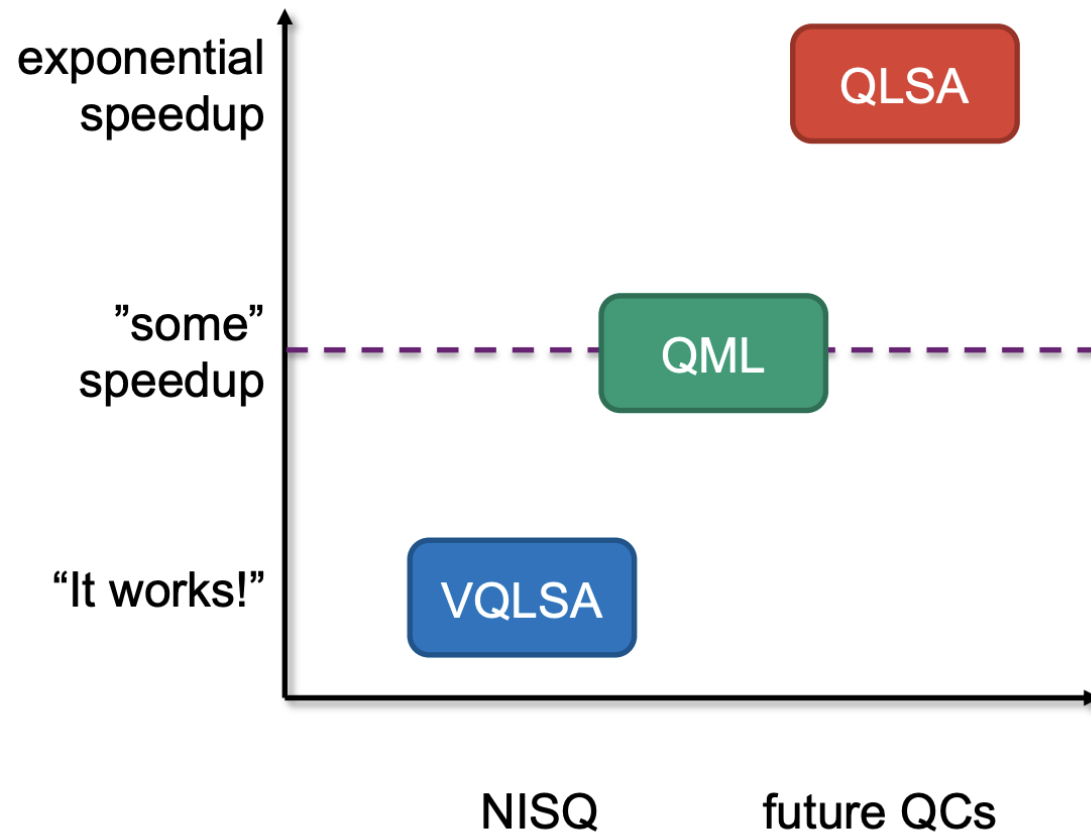
Harrow, A. W., Hassidim, A., & Lloyd, S. (2009). Quantum algorithm for linear systems of equations. *Physical review letters*, 103(15), 150502.

Is that enough?

- HHL (09) claims exponential speed-up $O(s^2 \kappa^2 \log(N)/\epsilon)$
- Textbook Gaussian elimination requires $O(N^3)$. Is this the best a classical solver can do?
- What's the cost of SPAM? State tomography will ruin all the speed up



What should we aim for?



At least a **quadratic speedup** is needed considering the clock speed difference.

Classical: GHz


Quantum: kHz-MHz

Timeline

- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)

Sparse linear system

- Most matrices arising from real applications are sparse



Example sparse A matrix
in CFD simulations


- Sparse matrix collection <https://sparse.tamu.edu>

Jargons

- **Sparsity:** number of elements are non-zero
- **Condition number:** measures how much the output value of the function can change for a small change in the input argument

$$\kappa(A) = \|A\| \cdot \|A^{-1}\|$$

- **Bandwidth:** the maximum number of diagonals on either side of the main diagonal that contain non-zero elements



Zoo

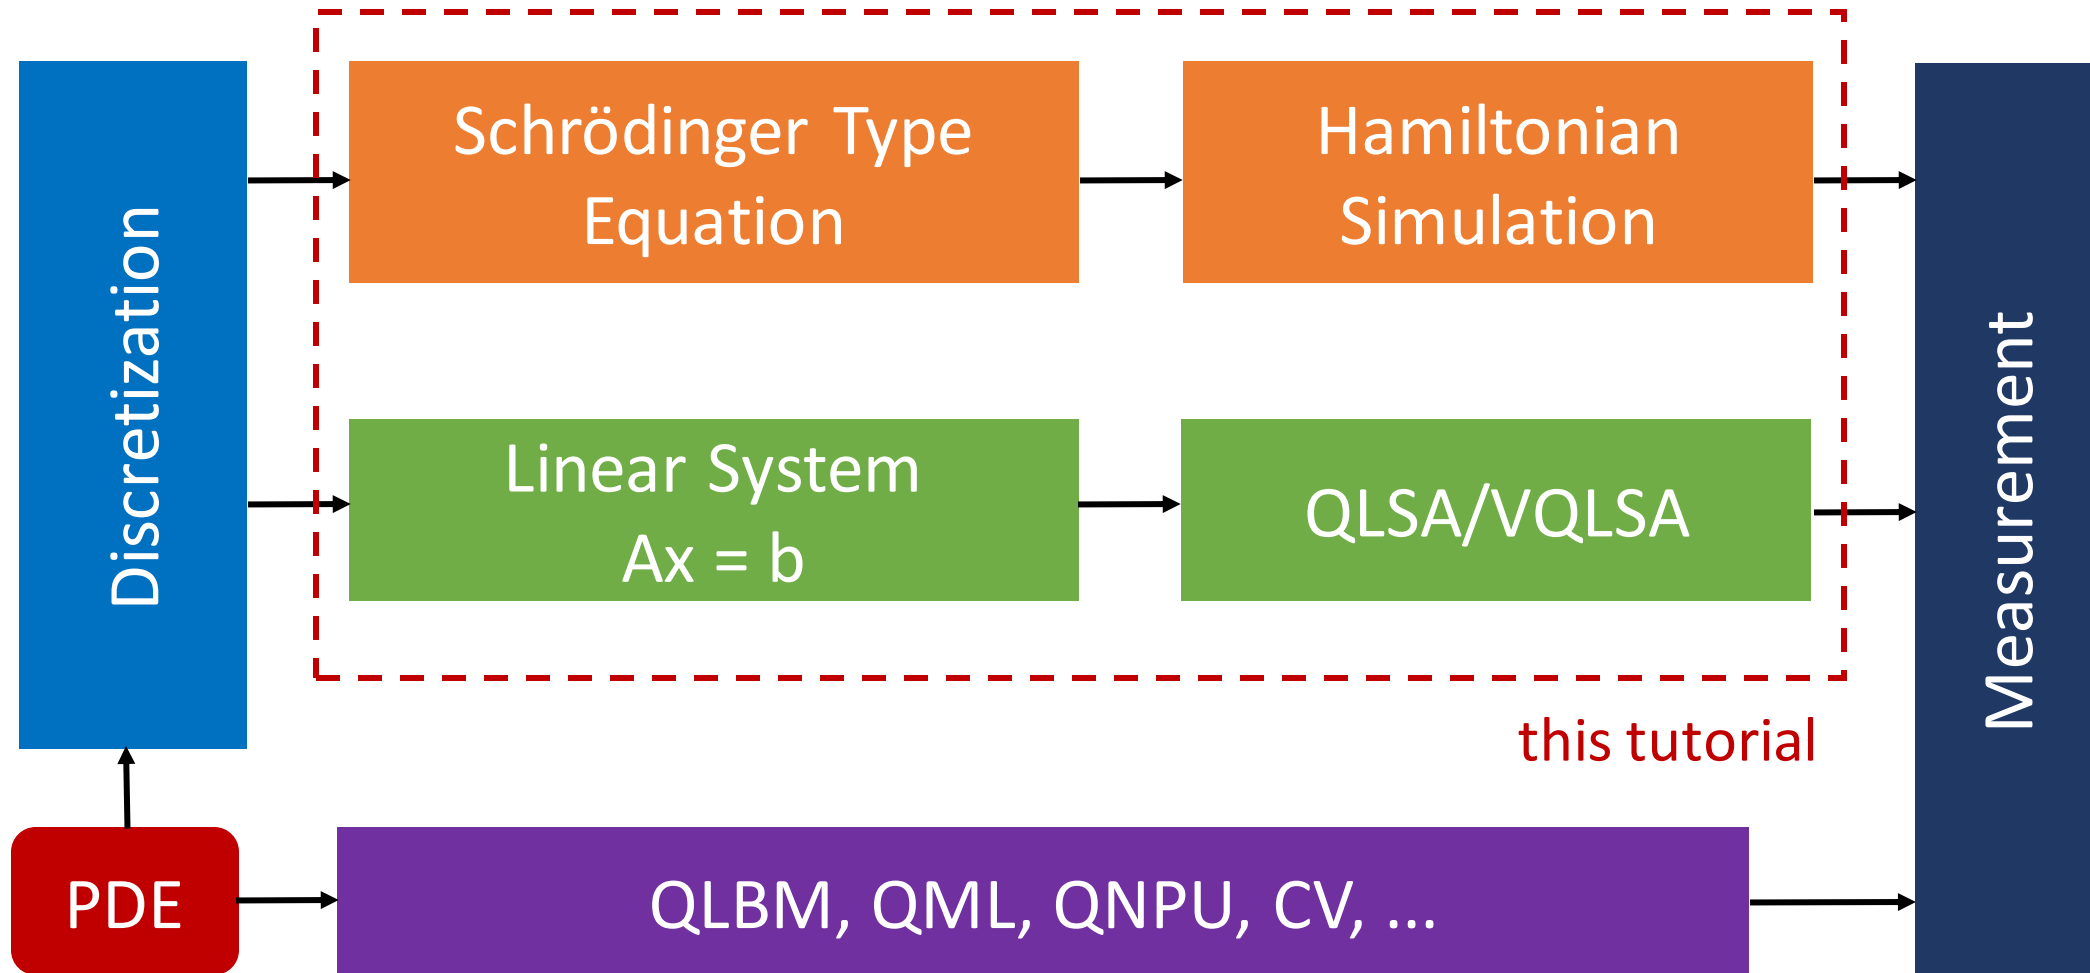
- Direct methods
 - LU decomposition (Gaussian Elimination) $O(N^3)$
- Iterative methods
 - Jacobi/Gauss-Seidel $O(N^2)$
 - Conjugate Gradient (CG) $O(N\sqrt{\kappa})$
 - Multigrid $O(N)$



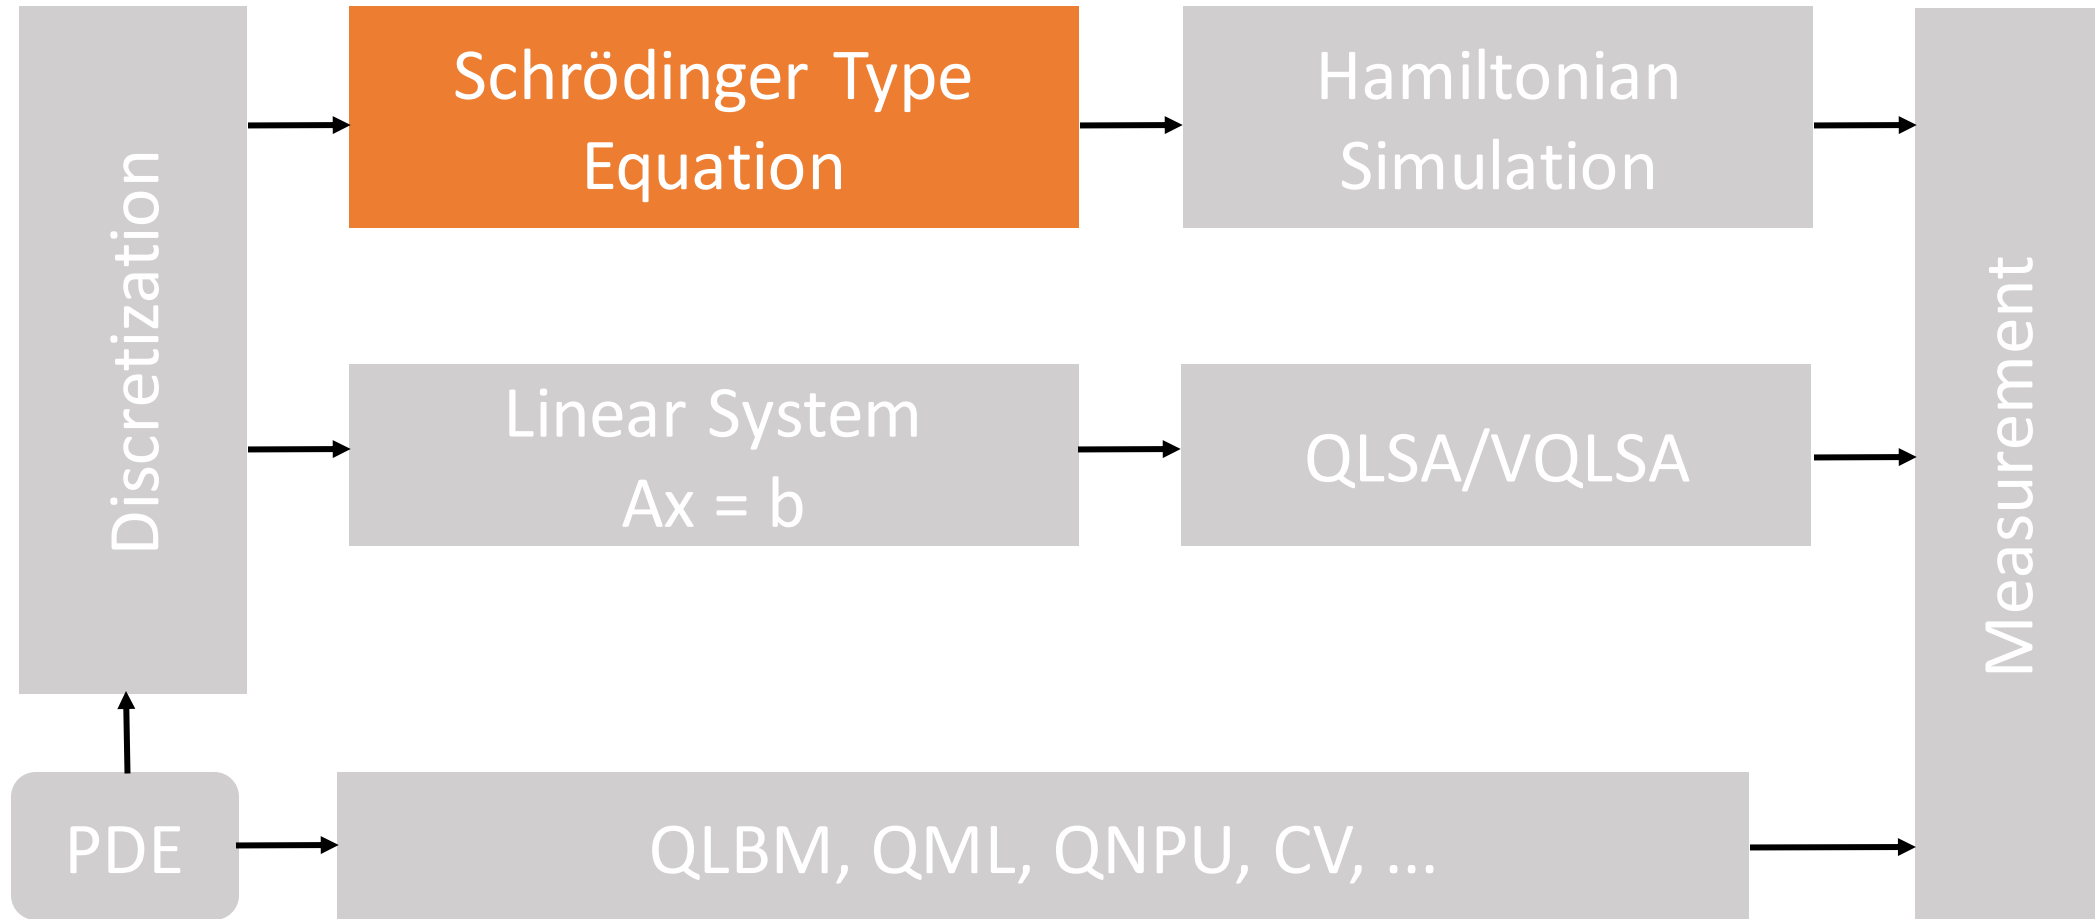
Timeline

- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)

Quantum Algorithms for solving PDEs



Quantum Algorithms for solving PDEs




Schrödinger Equation

- Goal is to solve $\frac{d\psi}{dt} = iH\psi$

- Can be achieved by efficiently implementing unitary operator

$$|\psi\rangle = e^{-iHt} |\psi_0\rangle$$

- The most natural type of PDE to solve on a quantum computer
- Some PDEs can be converted into this scheme



Wave Equation

- Goal: $\frac{\partial^2 f}{\partial t^2} = -\Delta f$
- Observe derivative of the Schrödinger's equation $\frac{d^2 \psi}{dt^2} = -H^2 \psi$
- Hermitian matrix

$$H = \begin{pmatrix} 0 & B \\ B^\dagger & 0 \end{pmatrix} \Rightarrow H^2 = \begin{pmatrix} BB^\dagger & 0 \\ 0 & B^\dagger B \end{pmatrix}$$

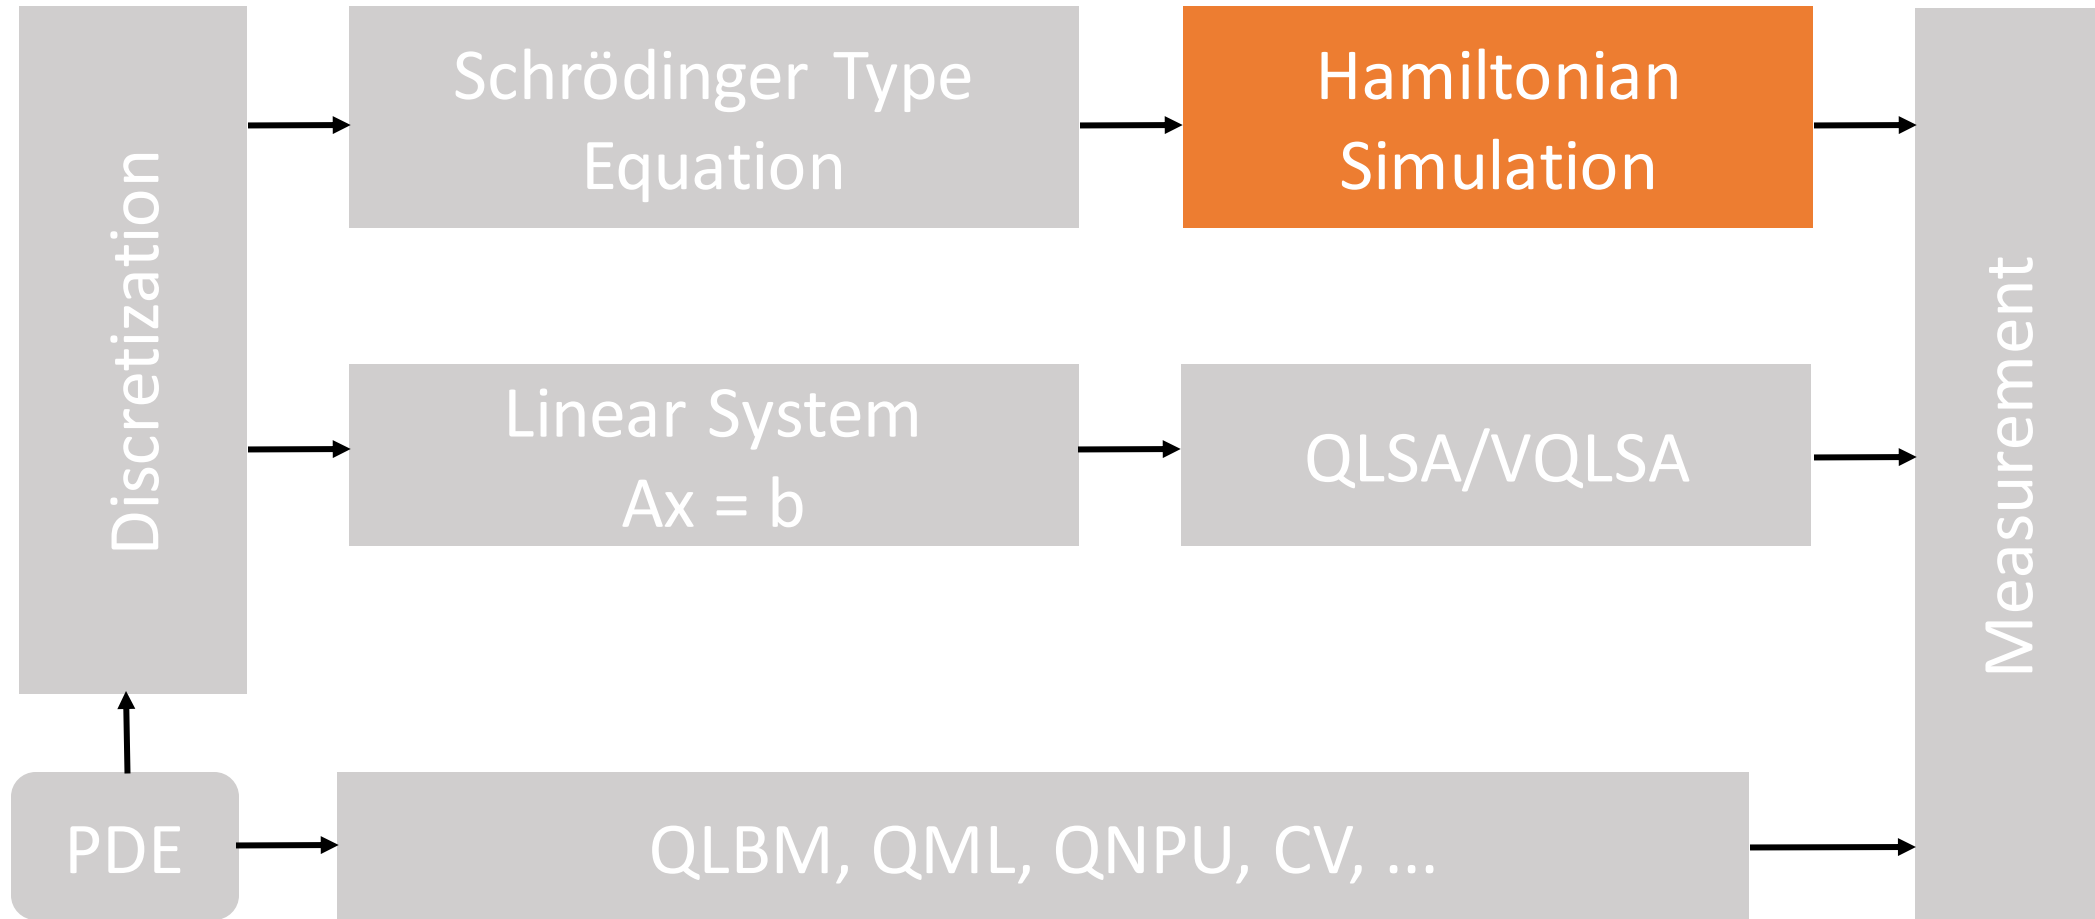
Laplacian

- Can be achieved by efficiently implement the unitary operator

$$|\psi\rangle = e^{-iHt} |\psi_0\rangle$$

- Not all PDEs can be written as a Schrödinger's equation!!


Quantum Algorithms for solving PDEs



Hamiltonian Simulation

Given: Hamiltonian

Goal: find algorithm to approximate $\|U - e^{iHt}\| \leq \epsilon$



- Decompose into local Hamiltonians


$$H = \sum_k H_k$$

- Suzuki-Trotter decomposition

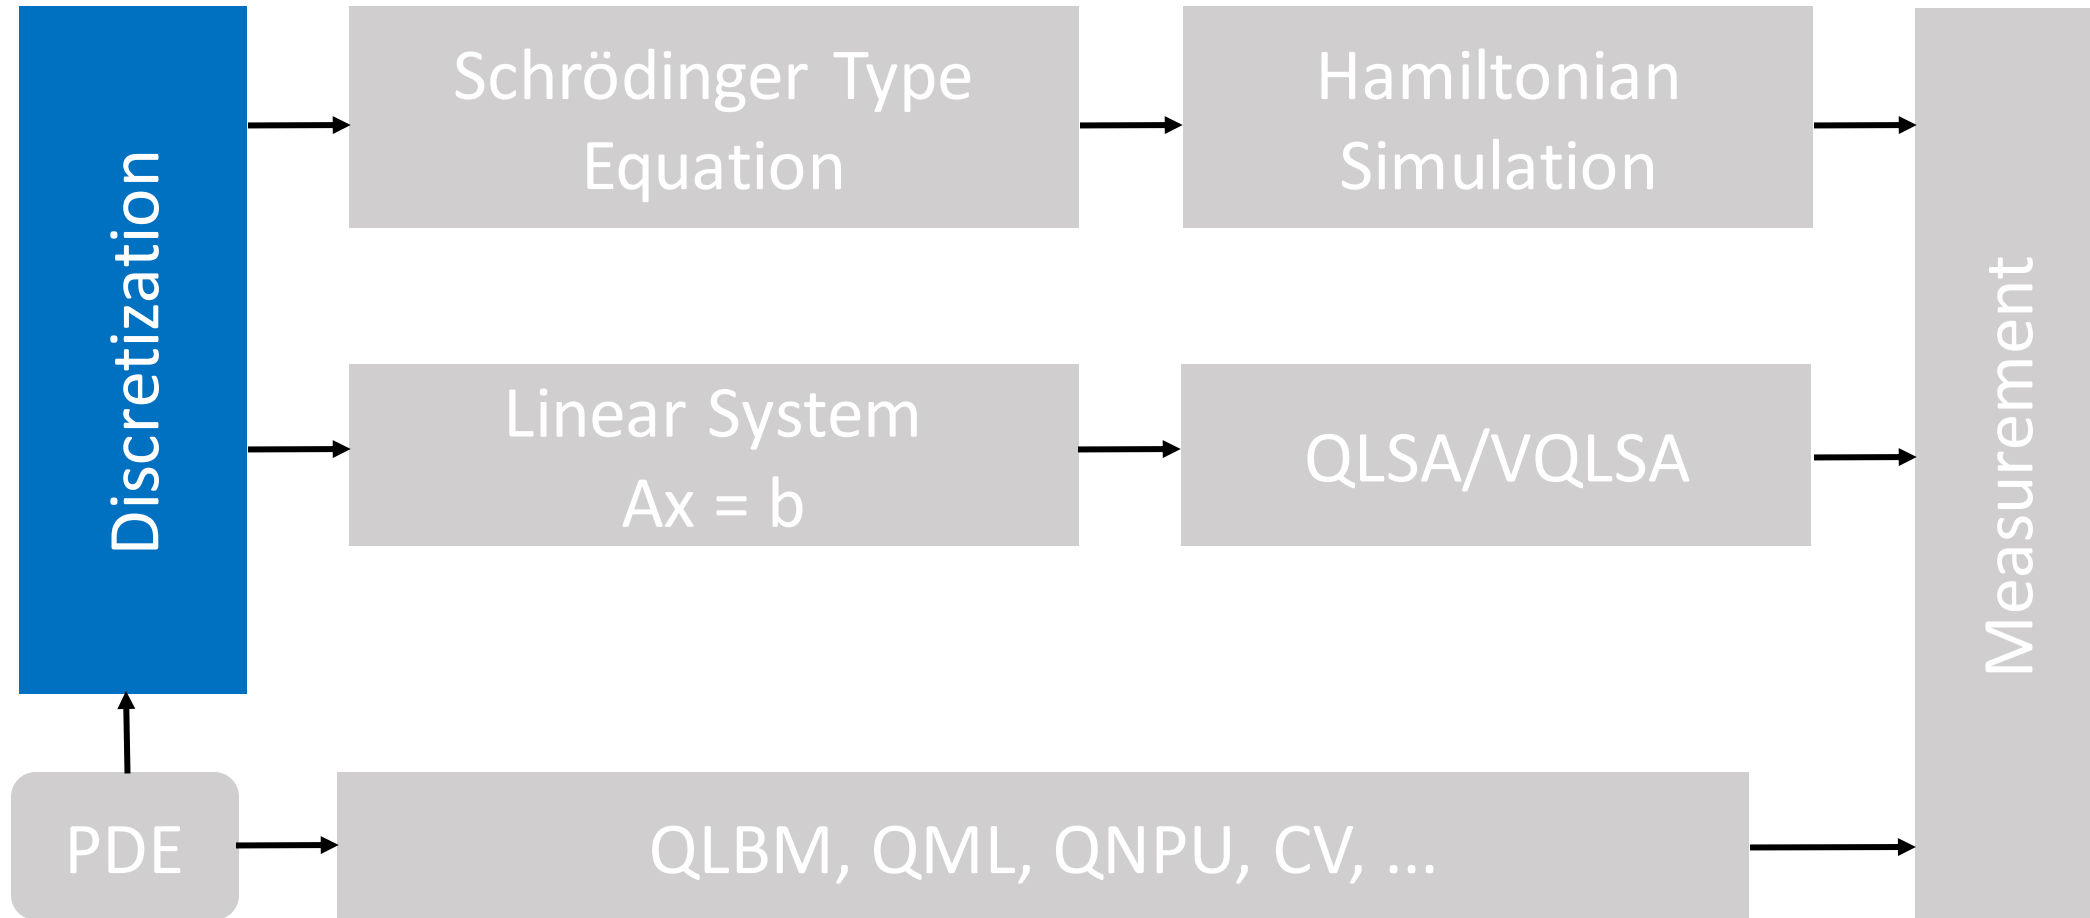
$$e^{-iHt} = (e^{-iH\tau})^r = \left(\prod_{k=1}^L e^{-iH_k\tau} + O(\tau^2) \right)^r = \left(\prod_{k=1}^L e^{-iH_k\tau} \right)^r + O\left(\frac{t^2}{r}\right)$$

Analog V.S. Digital

- The target Hamiltonian is close to the system Hamiltonian itself
- So far, most meaningful results are from analog simulators
- Digital simulation (gate based) is difficult due to hardware limitations



Quantum Algorithms for solving PDEs




Example

- Discretize the 1D spatial domain into grid of n points and approximate derivative via central difference (CDS)

$$\frac{\partial^2 \phi(x)}{\partial x^2} \approx \frac{\phi_{i-1} - 2\phi_i + \phi_{i+1}}{h^2}$$

- 2D Laplacian

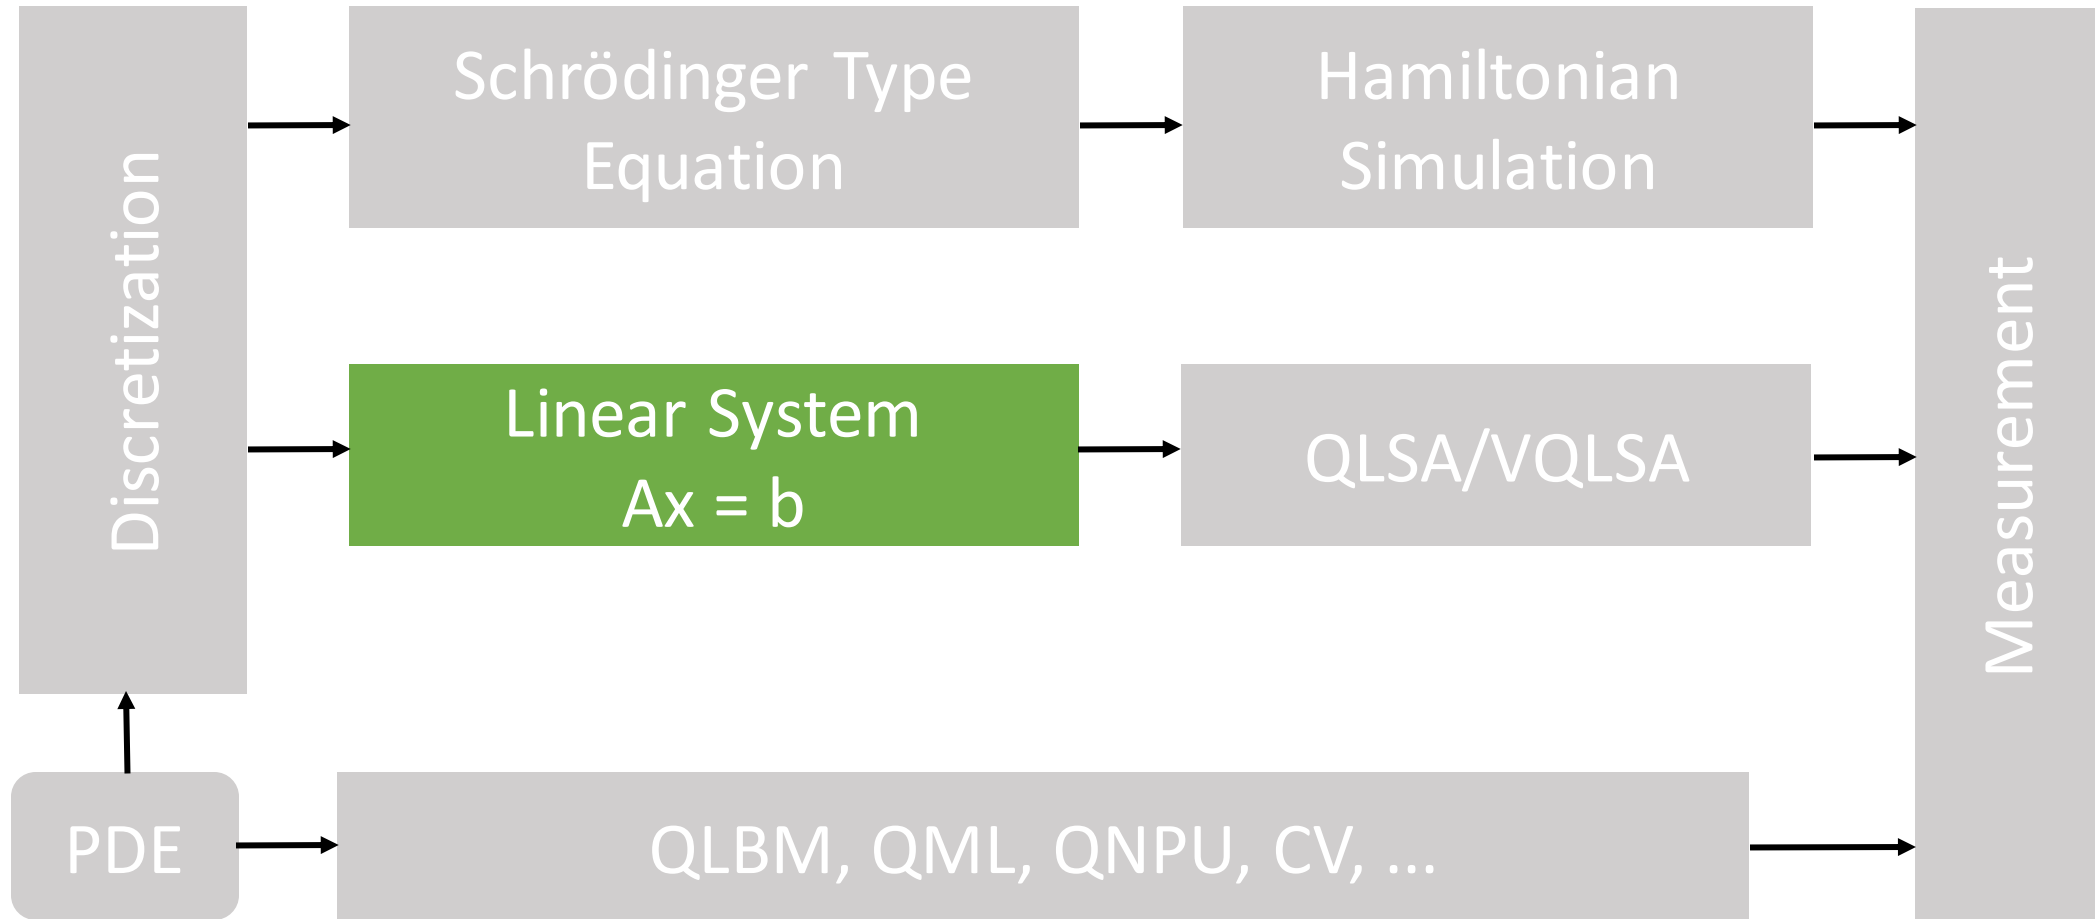
$$\nabla^2 \phi_{i,j} = \frac{1}{h^2} (\phi_{i-1,j} + \phi_{i+1,j} + \phi_{i,j-1} + \phi_{i,j+1} - 4\phi_{i,j})$$



- For many equations, achieving a discretization error less than ϵ requires

$$N = \mathcal{O} \left(\text{poly} \left(\frac{1}{\epsilon^d} \right) \right)$$


Quantum Algorithms for solving PDEs



Linear system $Ax=b$

- For example, let's consider a 1D Poisson equation

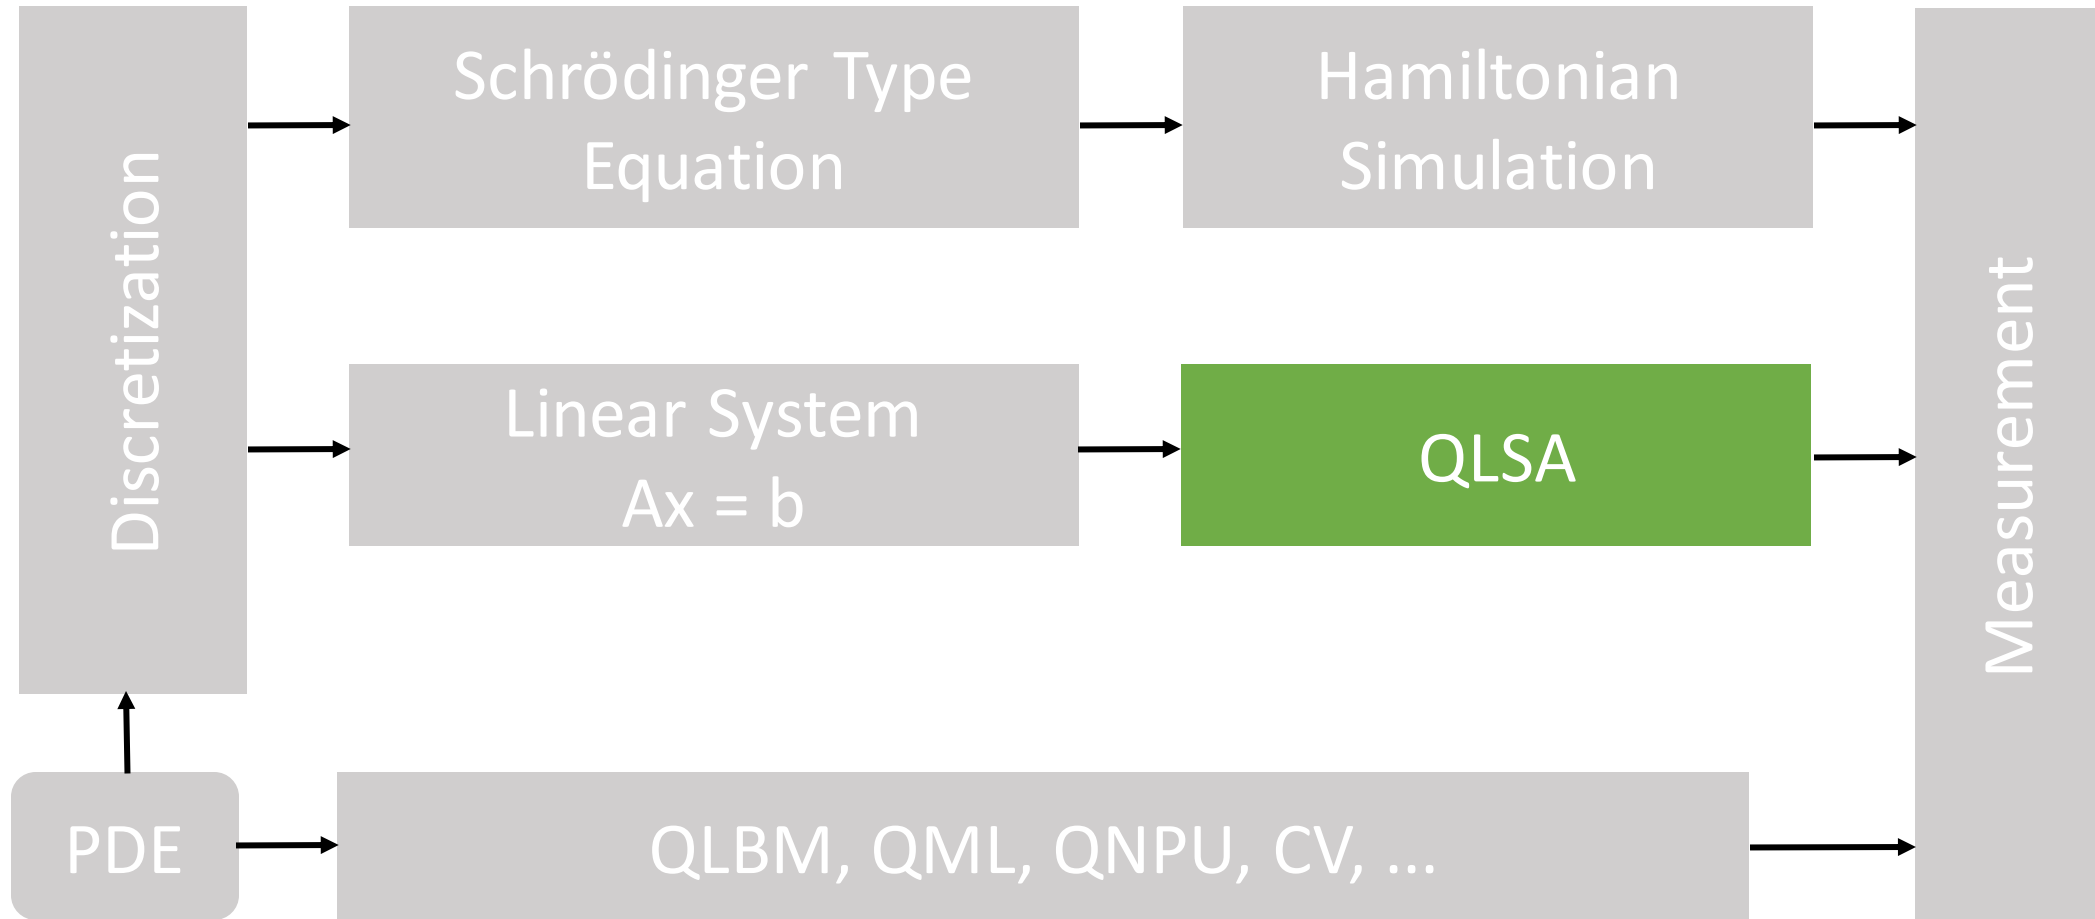
$$\Delta\phi(x) = f(x), \quad \text{with} \quad \phi(a) = \phi_a, \phi(b) = \phi_b.$$



- Apply the central difference scheme (CDS), we get

$$A = \begin{bmatrix} -2 & 1 & 0 & \cdots & 0 & 0 \\ 1 & -2 & 1 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \cdots & 1 & -2 \end{bmatrix}, b = \begin{bmatrix} h^2 f_1 - \phi_a \\ h^2 f_2 \\ \vdots \\ h^2 f_{n-1} \\ h^2 f_n - \phi_b \end{bmatrix}.$$

Quantum Algorithms for solving PDEs

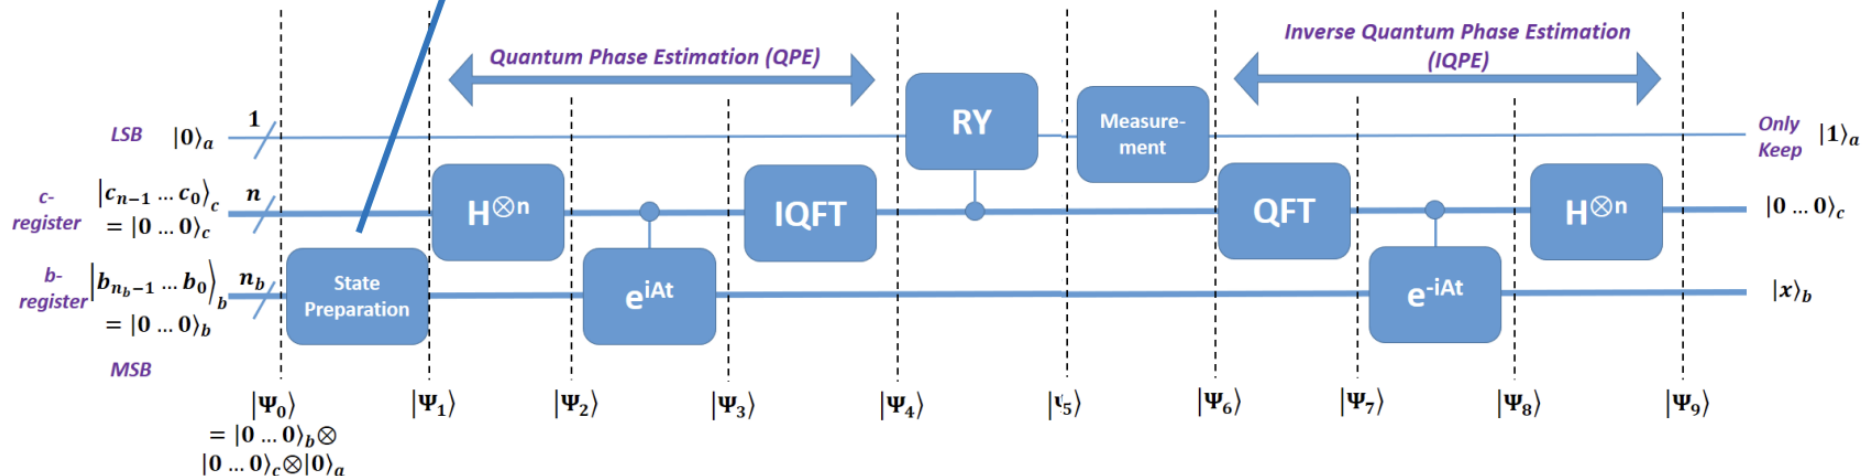


HHL Algorithm

$$O(s^2 \kappa^2 \log(N)/\epsilon)$$

- Goal: prepare the quantum state $|x\rangle$ that solves $Ax = b$

$$A = \sum_{j=0}^{2^n-1} \lambda_j |u_j\rangle\langle u_j|, \quad b = \sum_{j=0}^{2^n-1} b_j |u_j\rangle, \quad |x\rangle = A^{-1}|b\rangle = \sum_{j=0}^{2^n-1} \lambda_j^{-1} b_j |u_j\rangle$$



State preparation


- General states cannot be prepared efficiently, not even approximated!
- Requires $O(N)$ gates using uniformly controlled rotations
- Certain states of the form $|\psi\rangle = \sum_i \sqrt{p_i} |i\rangle$ can be prepared efficiently, e.g., using quantum GANs with $\mathcal{O}(\text{polylog}(N))$ gates

$$\frac{\partial^2 \log(p(x))}{\partial x^2} < 0 \quad \text{such as normal distribution}$$

- Possible to reduce time complexity by adding ancillary qubits

Measurement

- Access to full quantum solution encoded on amplitudes require Quantum State Tomography (QST)
- This requires $\mathcal{O}(4^n)$ of state sampling + MLE post-processing
- Cost few hours to reconstruct a ~ 20 qubit density matrix and this kill the exponential speed-up!



Measurement

- The global solution cannot be easily extracted in general
- Many functions of a state that can be extracted efficiently
 - Inner product estimation (SWAP test)
 - Small set of amplitudes (amplitude estimation)
- Measure any functional requires $\mathcal{O}(\text{poly}(1/\epsilon))$ repetitions
- This cost is unavoidable due to the nature of quantum computing!

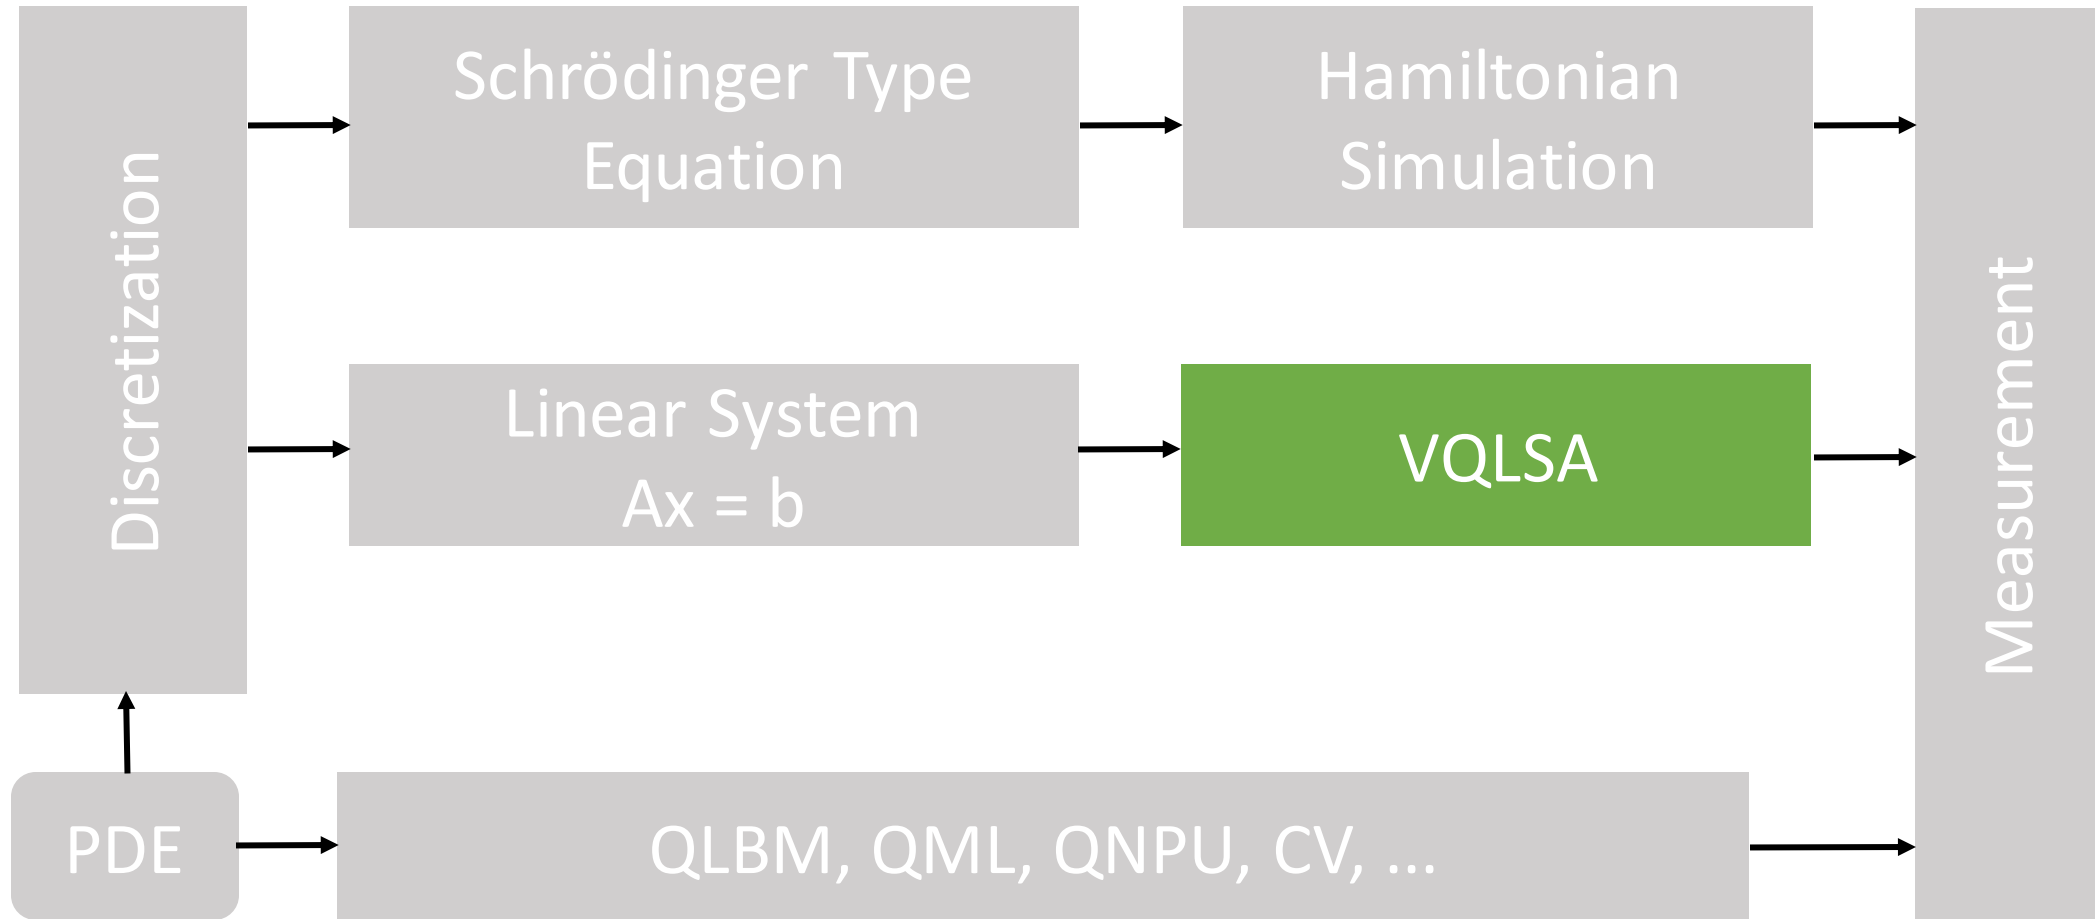
Timeline

- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)

Timeline

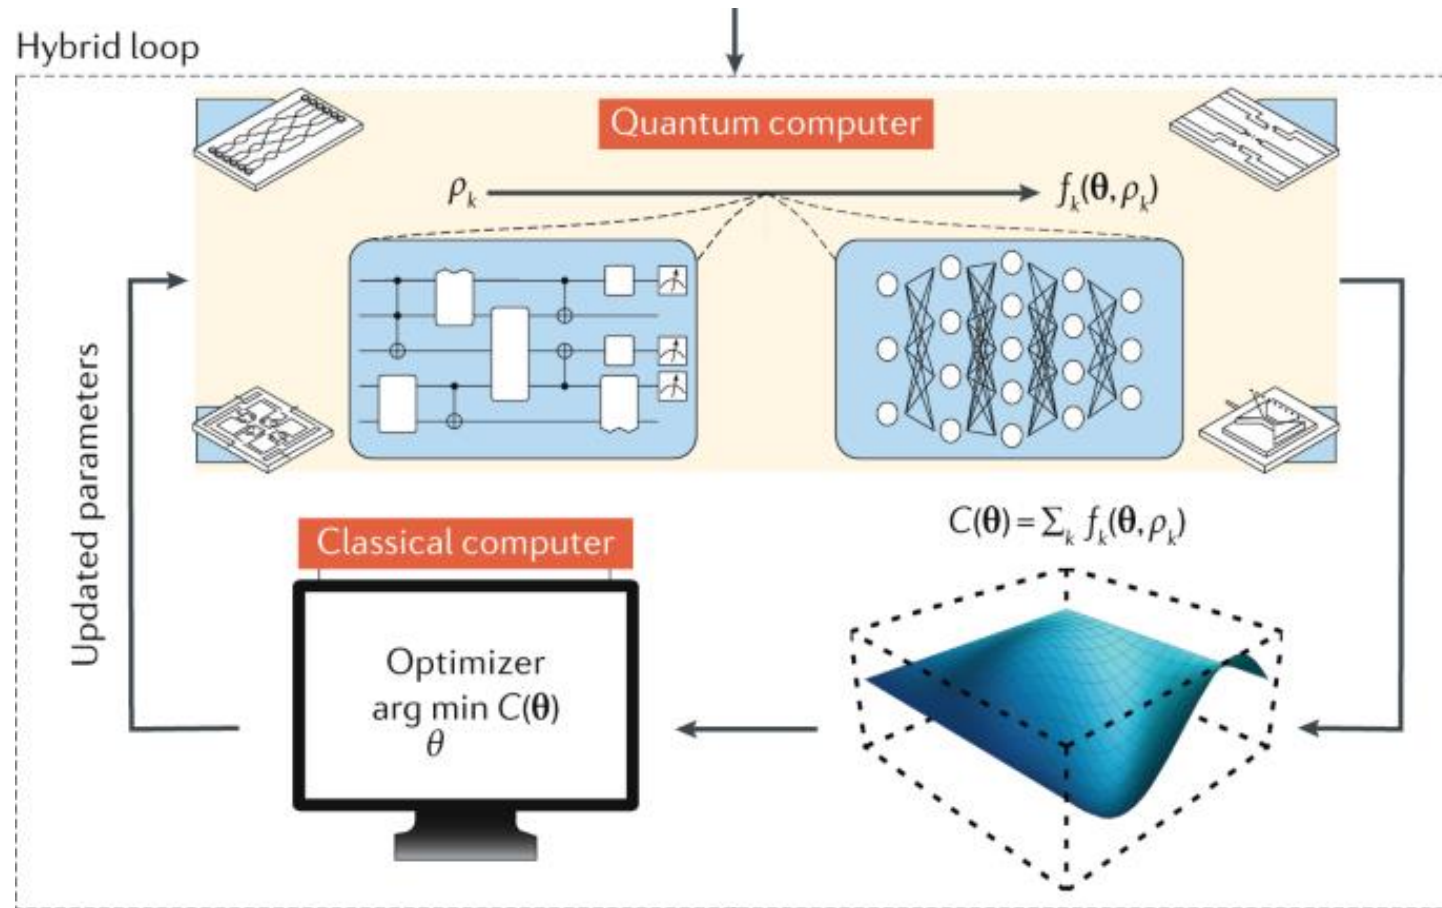
- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)

Quantum Algorithms for solving PDEs

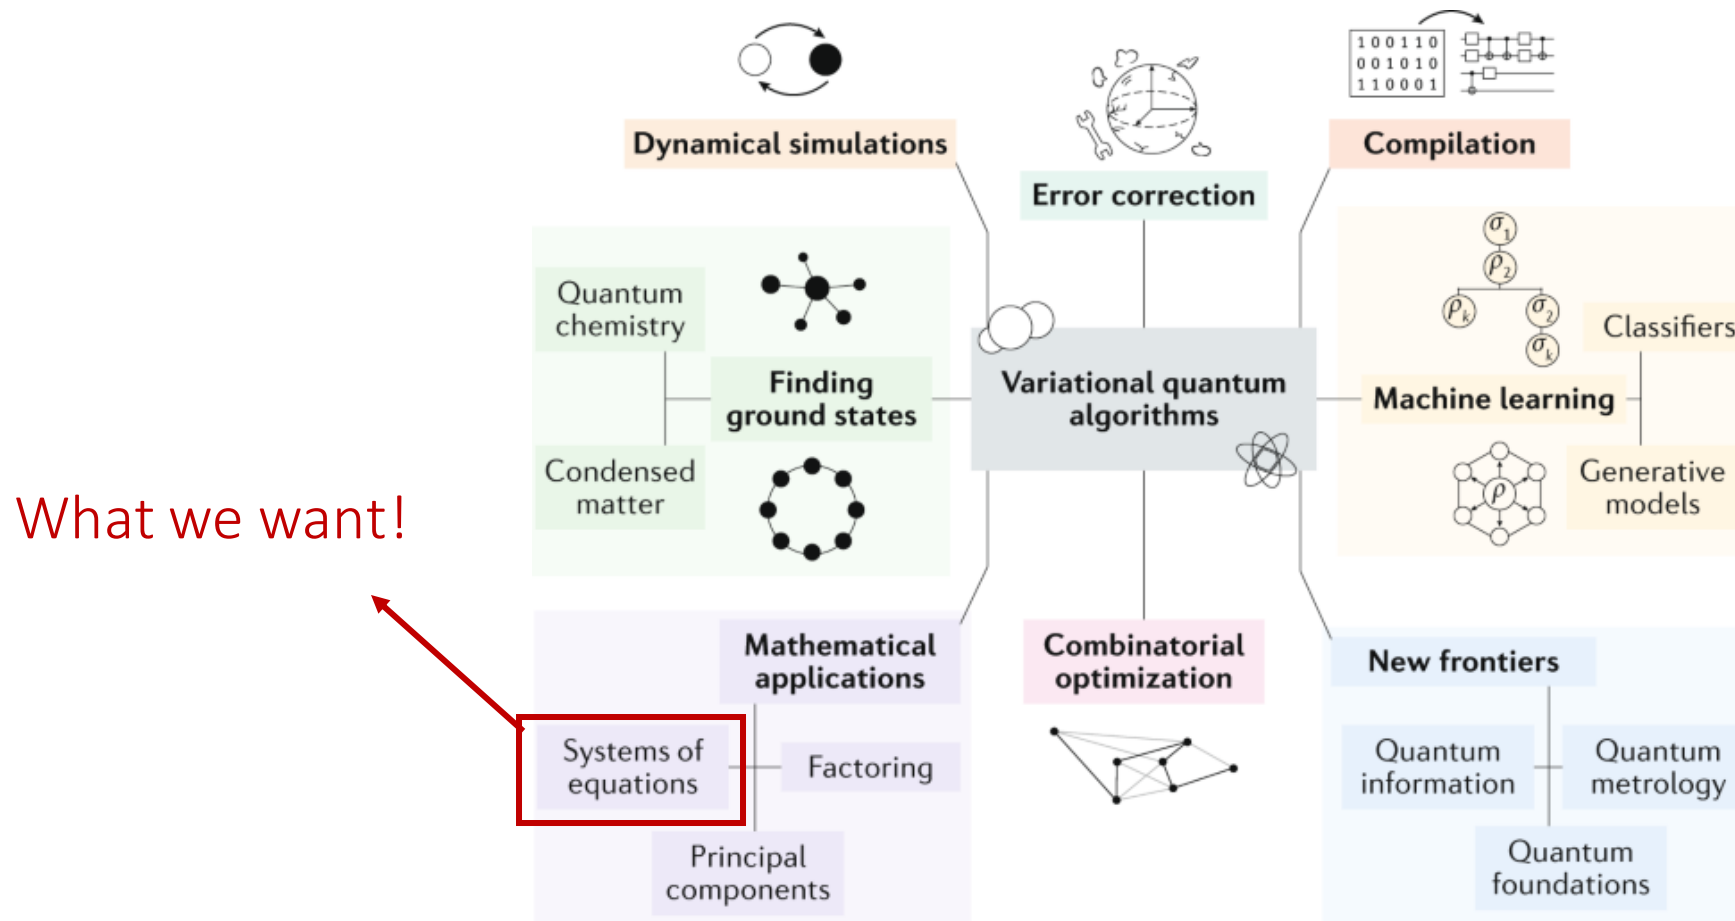


Near-term QC

In most cases, people are referring to some Variational Quantum Algorithms (VQAs)

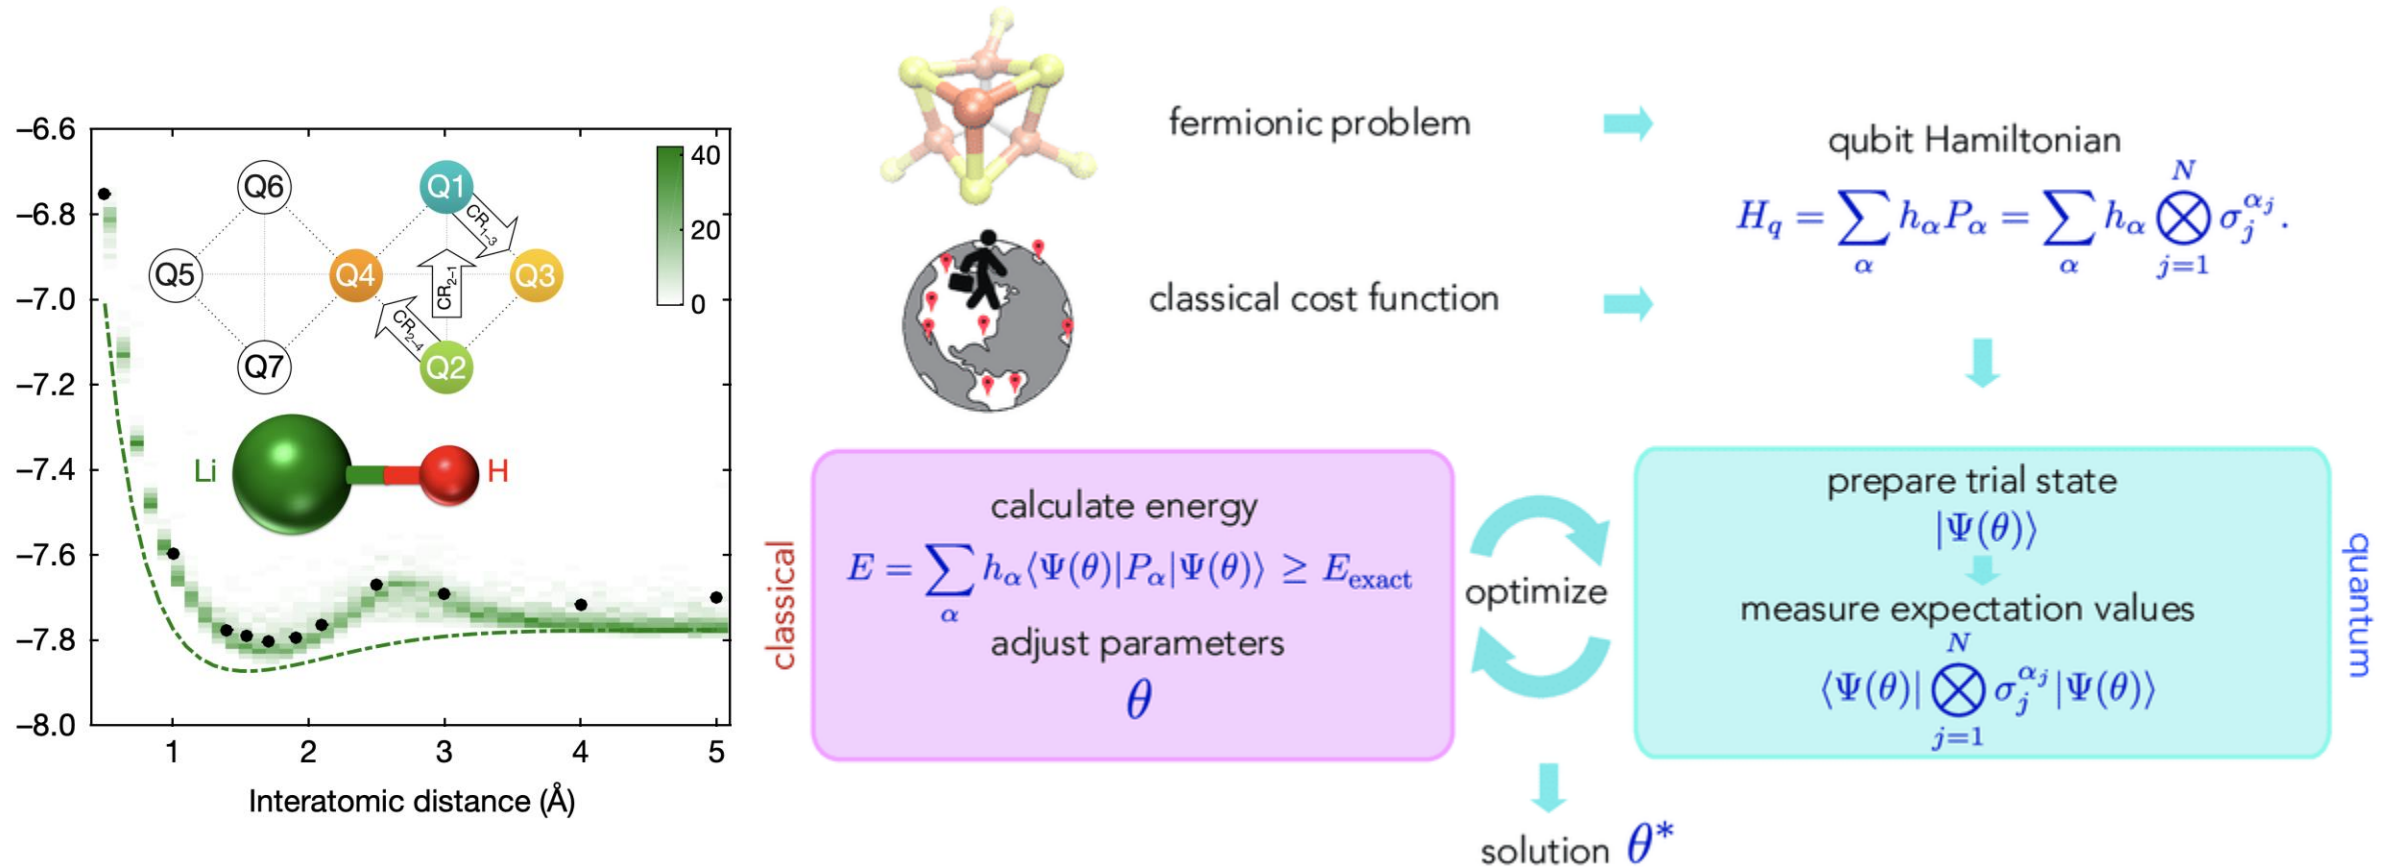


Variational quantum algorithms



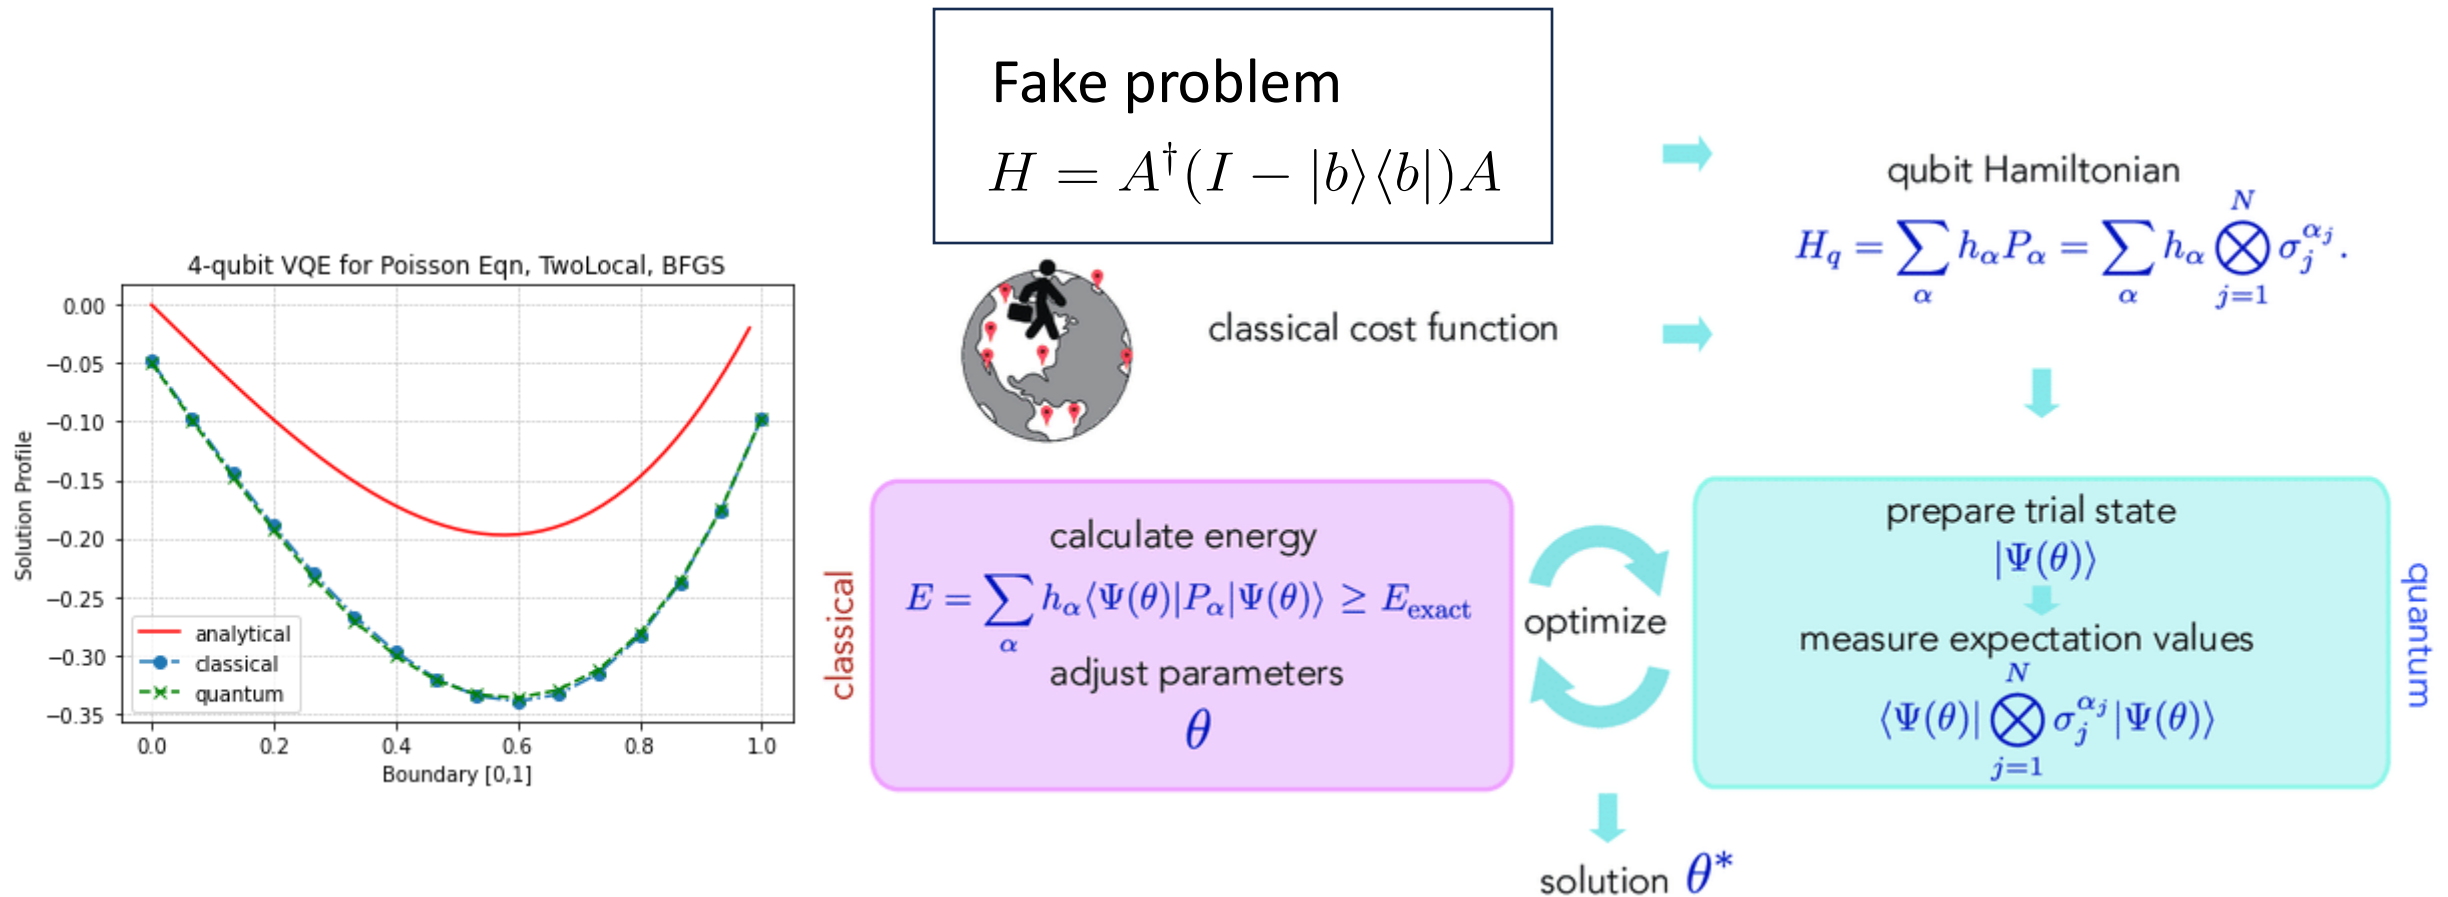
Variational Quantum Eigensolver

VQE adopts the variational principle to estimate the ground state energy of molecules



VQE as VQLSA

Encode the solution of linear system into the ground state of a target Hamiltonian

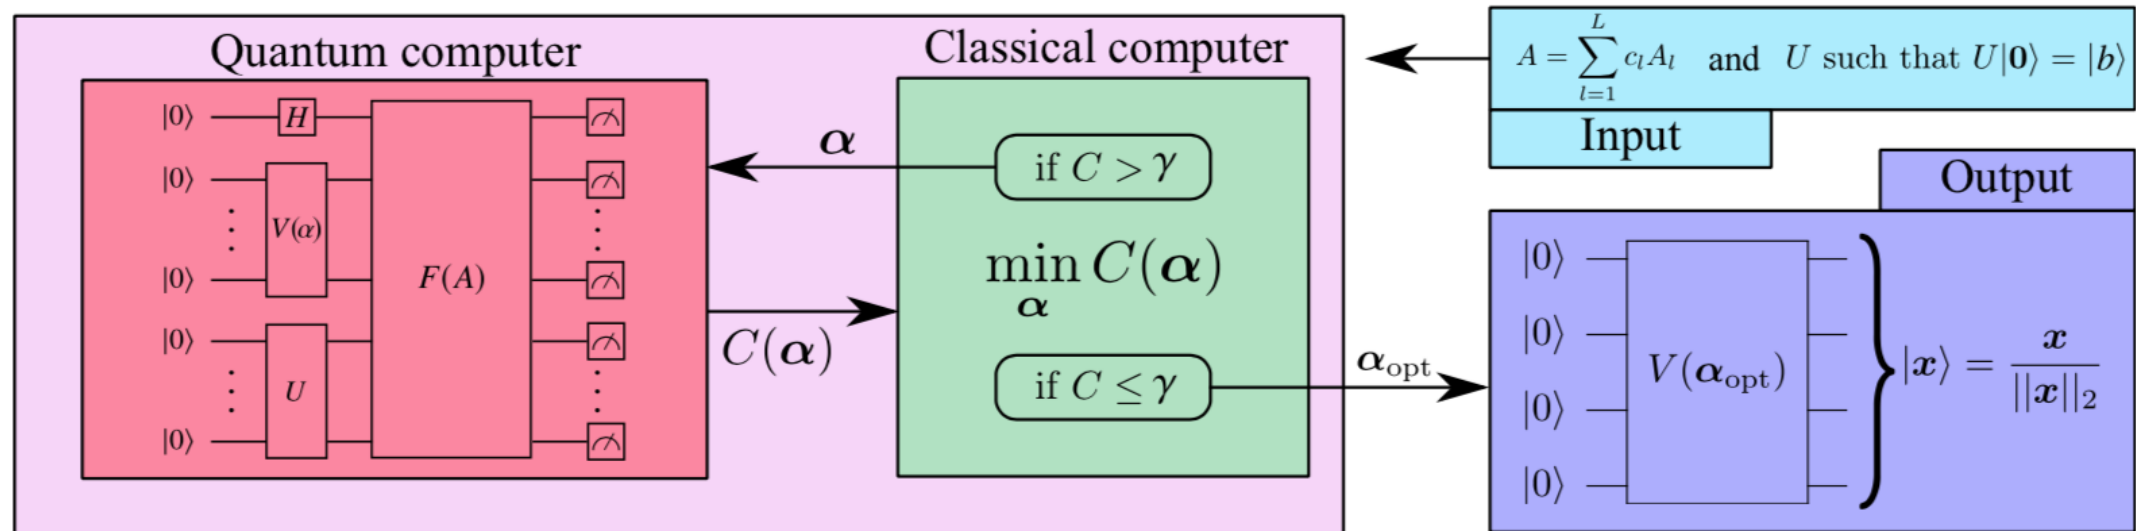


Variational Quantum Linear Solver

Similar idea but different cost function:

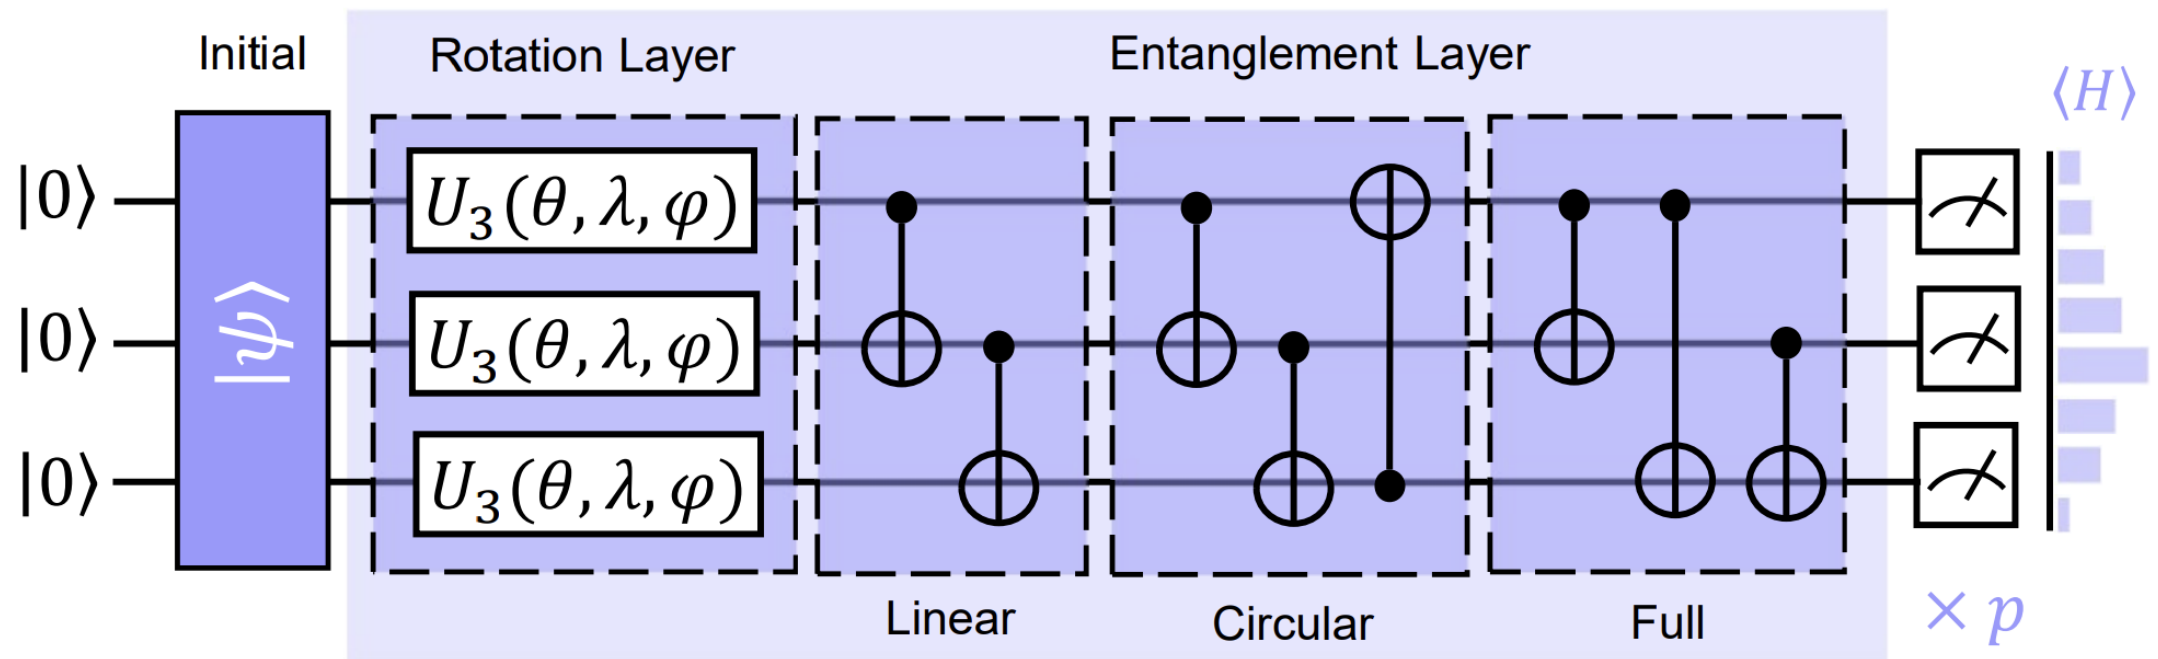
$$\left. \begin{array}{l} |\Phi\rangle \perp |b\rangle \Rightarrow C(\Theta) \text{ large} \\ |\Phi\rangle // |b\rangle \Rightarrow C(\Theta) \text{ small} \end{array} \right\} |\Phi\rangle = A|\psi(\Theta)\rangle$$

$$\hat{C}(\Theta) = 1 - \frac{|\langle \Phi | b \rangle|^2}{\langle \Phi | \Phi \rangle}$$



Ansatz

- Find optimal problem-specific circuit structure is difficult due to large search space
- Hardware-Efficient Ansatz (HEA) is usually the choice if we know nothing about the problem structure or symmetry




Optimizer

- Gradient based (converge faster)
 - Gradient Descent (GD)
 - SGD
 - Adam
 - BFGS (my favorite)
- Non-gradient based (more robust to noise)
 - COBYLA
 - Nelder-Mead
 - Powell
 - SPSA

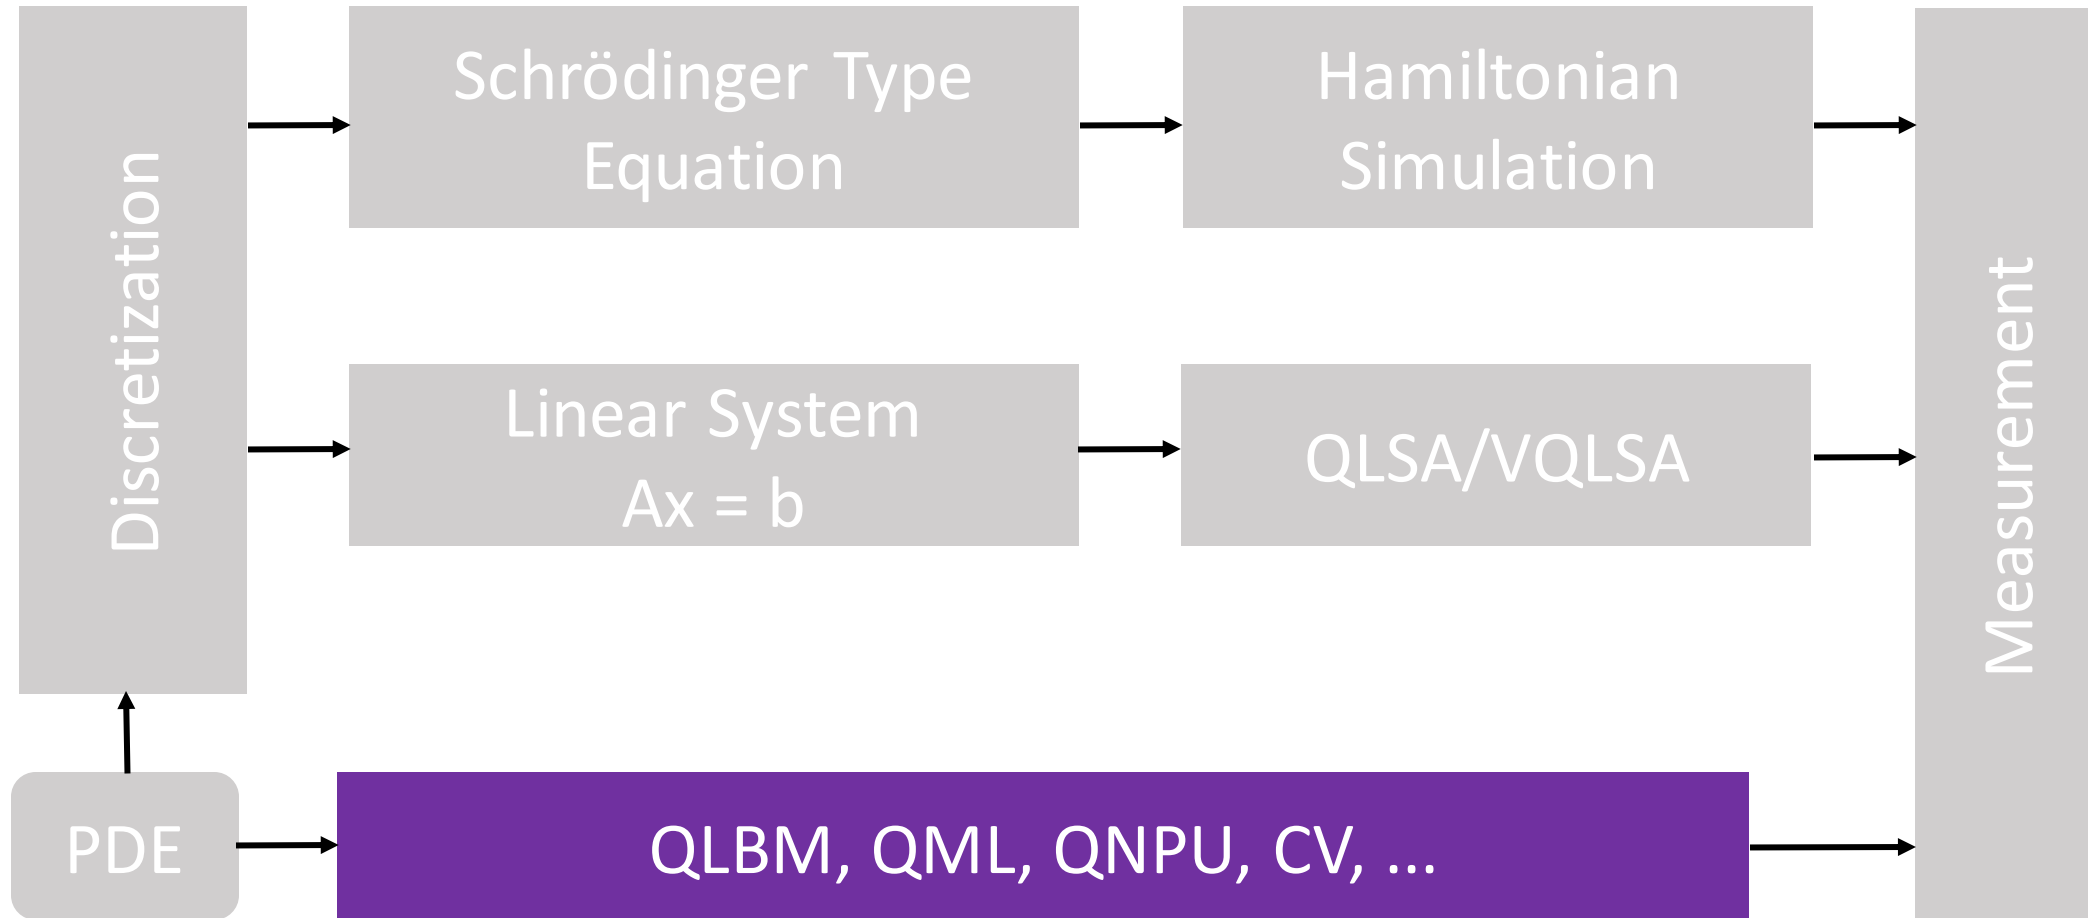
There is no best choice. You need to test various things.

Barren Plateau (BP)

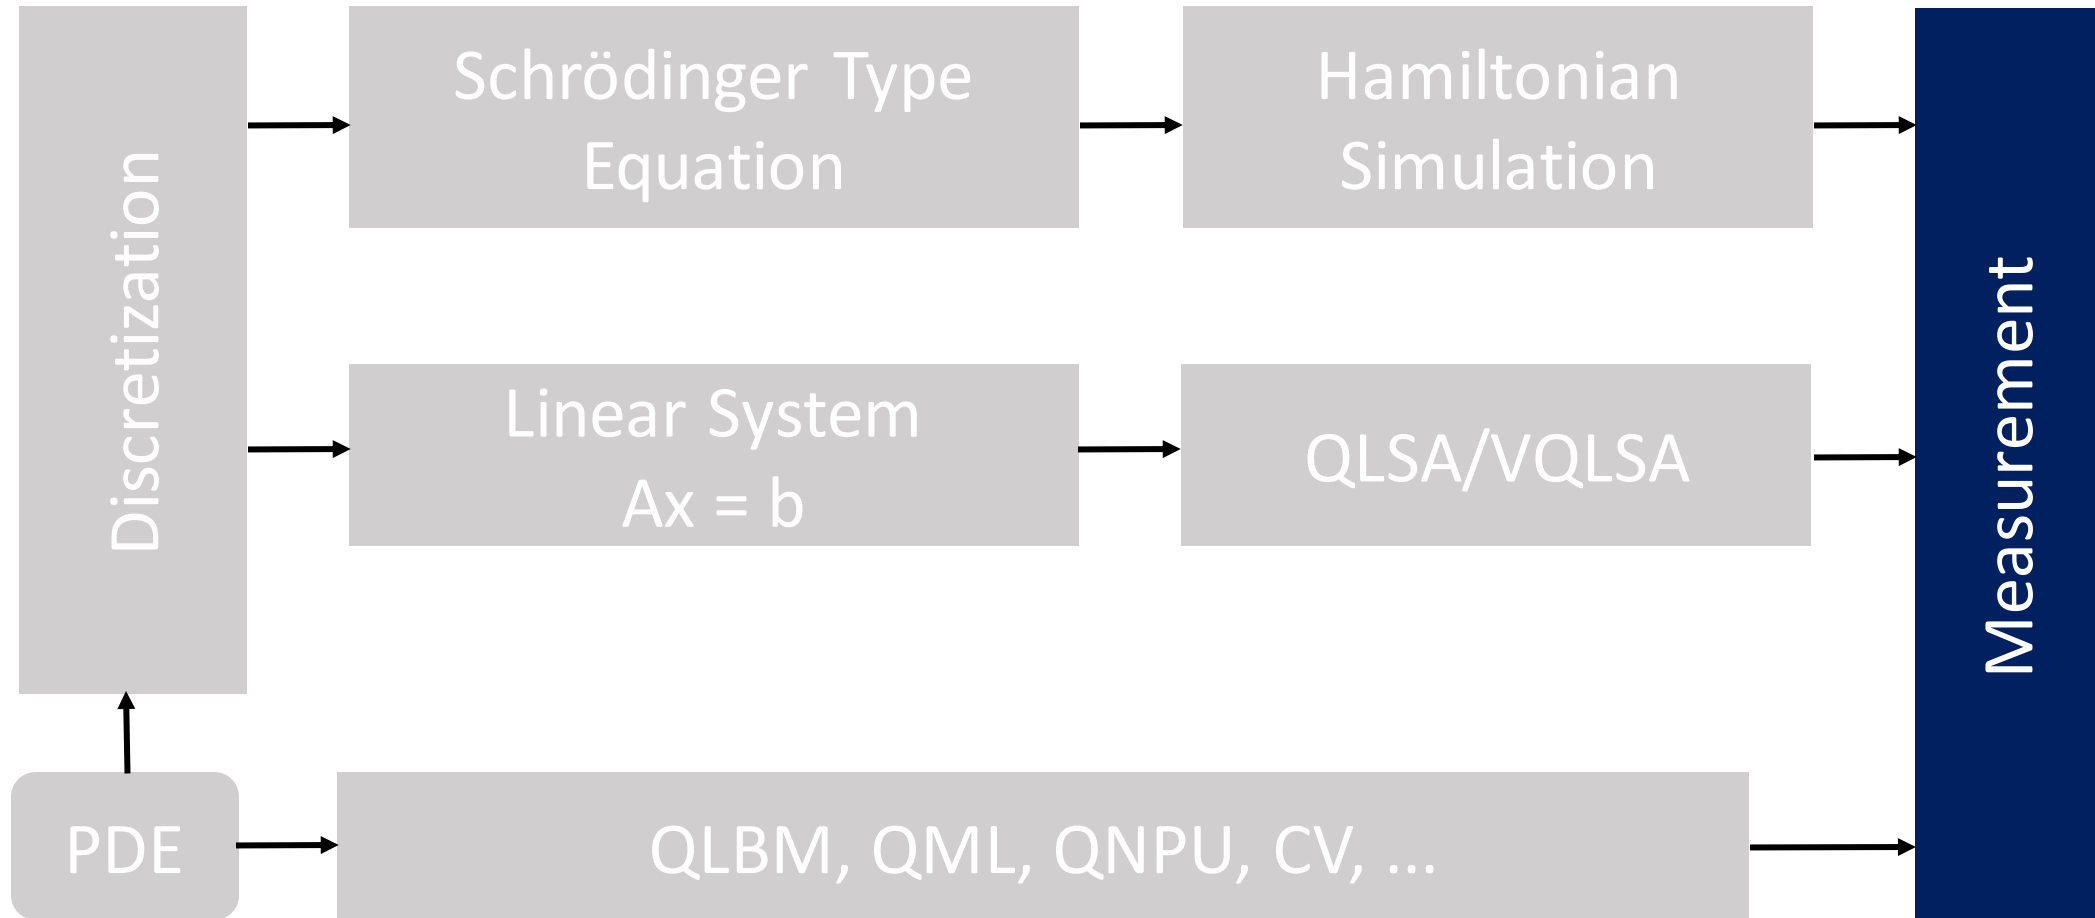
- When a given cost function exhibits a BP, the magnitude of its partial derivatives will be, on average, exponentially vanishing with the system size
- This makes VQE+HEA training difficult for large qubits



Quantum Algorithms for solving PDEs



Quantum Algorithms for solving PDEs




Fidelity Estimation

- Expand two states in the Pauli basis as

$$\rho = \sum_j a_j W_j / 2^{N/2} \quad \psi = \sum_j b_j W_j / 2^{N/2}$$

- The fidelity can be estimated as $F(\psi, \rho) = \text{tr}(\psi\rho) = \sum_j a_j b_j = \sum_j \left(\frac{a_j}{b_j} \right) b_j^2$
- Notice $a_j/b_j = \text{tr}(W_j\rho) / \text{tr}(W_j\psi)$
- Then we can sample these observables to obtain a randomized estimation of fidelity with sample $O(2^n/\epsilon^2)$



Take-home message

- Steady state problems are more suitable to solve for near-term
- Quantum algorithms are not good at nonlinearity
- Explore more near-term algorithms for linear system
- Possible research direction is simulating QLSA such as HHL and improved version in HPC systems

Timeline

- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)

Thanks for your attention!

- Q&A
- Coding time!
- We are looking for collaborations if you are interested.

Timeline


- Session 1 (10-11:30 am)
 - Introduction (20 min)
 - Classical Solution (10 min)
 - Quantum Solution (30 min)
 - Notebook on Hamiltonian Simulation (30 min)
- Session 2 (1-2:30 pm)
 - Near-term QC (30 min)
 - Notebook on VQA (40 min)
 - Open discussion (20 min)

Software

- QuDiffEq.jl <https://github.com/QuantumBFS/QuDiffEq.jl>
- SimuQ <https://pickspeng.github.io/SimuQ/>

How to choose your backend?

- Through word of mouth, many people are in favor of Quantinuum's hardware due to high fidelity.
- What's your experience on this?
- Have you considered to run your algorithm on the actual quantum hardware?



How to access backend?

- OLCF
- AWS Research Credit
- Microsoft Azure Credit
- IonQ Research Proposal
- IBM Quantum free/paid plans